

Transfer Learning Results

This is a recurring report of the transfer learning experiment, including **TCA, JDA, Baseline: ResNet50, Baseline: Digit Network, Self-ensemble Visual Domain Adapt Master, DANN, ADA, MCD_UDA, MADA.**

The above line is just to support the display of mathematical formulas(may be it doesn't appear), it is recommended to open with **Chrome**.

[Data set](#) (downloaded from all over, and integrate by my. The extraction code: zhgy)

黄晨晰 Chenxi Huang

Computer and science majors

Email: mrsemppress98@gmail.com / hcx_98@163.com

Web: <http://mrsemppress.top>

Transfer Learning Results

Preface

Non-deep transfer learning

TCA

Compare

JDA

Refer to Long's

Compare

Refer to Wang's

Compare

Compare with TCA

Deep transfer learning

ResNet50

Compare

With Wasserstein

Digit Network

Compare

Self-ensemble Visual Domain Adapt Master

Compare

Domain Adversarial Training of Neural Networks

Compare

ADA

Compare

MCD_UDA

Svhn to mnist

Three layer fully connected network
using gradient reversal layer

Mnist to usps

Three layer fully connected network
Four layer fully connected network

Compare

MADA

Compare

Post-preface

Preface

The machine parameters used are as follows:

CPU	内存	显卡
i7-5930K (6/12)	125G	GTX 1080 * 4

I also read codes of other algorithm. And some codes try them myself, others are not.

Each one gives ways to execution, as well as attention. At the same time, while giving the results, the analysis was carried out, including the comparison with the previous algorithm, paper, and the analysis of different data sets (You can see the part of [Compare](#)).

Attention:

1. **[Record]** Due to space limitations, the experimental group released by each algorithm is limited. If you want to see more results of each group, see the `./results_image/*.png` .
2. **[Use log]** Since the way the data was previously recorded is directly output, then the screenshot shows. But the later the algorithm, the larger the amount of data, the more iterations, so it is not convenient for screenshots. I want to record by log, but it takes time to retest it. So I only give the code, if you want to get the data, please run it again. And you can find the log in `./results/model_name/source_to_target/model_name.csv` .
3. **[Input]** If you still have problems with the input, you can open the `main.py` file and use a basic example on each model.
4. **[Gpu]** Note that the default value of my gpu here is `3`, if you use the machine GPU does not have 4 blocks, you need to specify when entering the command.

Non-deep transfer learning

TCA

Download the data set of Office-31, the input format is `.mat` (if it is not, please use Matlab to convert it).

You can see the result with the command: `python TCA.py`:

Since I integrated the code, please use `python main.py --model='TCA'` and you can set other parameters.

The results is as follows:

(The accuracy is up to 91.08%)

Webcam to Dslr

```
[root@71b4a3d0c1b1:/home/huangchenxi# python TCA.py
TCA:  data = Office-31 lambda = 1 src: data/Office-31/webcam.mat tar: data/Office-31/dslr.mat
0.910828025477707
root@71b4a3d0c1b1:/home/huangchenxi# ]
```

After the unified framework, the test by default (changed a set of source and target domains, but the results are still in the normal range):

Amazon to Webcam

```
root at 60a5497e4409 in /home/huangchenxi/transfer-learning/transfer-learning
$ python main.py
TCA:  data = Office-31 lambda = 1 src: data/Office-31/amazon.mat tar: data/Office-31/webcam.mat
0.4
```

Compare

Since the paper's experiments were based on `Cross-domain WiFi Localization` and `Cross-domain Text Classification`, so I used another paper "Transfer Feature Learning with Joint Distribution Adaptation" to verify:

(Because I **add fine-tune**, so the result will **be better than paper**. You can see more detailed information in `/results_image/TCA-x-y.png`.)

Dataset	In paper(%)	In my test (%)
C->A	44.47	45.62
C->W	34.24	39.32
C->D	43.31	45.86
A->C	37.58	42.03
A->W	33.90	40.00
A->D	26.11	35.67
W->C	29.83	31.52
W->A	30.27	30.48
W->D	87.26	91.08
D->C	28.50	32.95
D->A	27.56	32.78
D->W	85.42	87.46

JDA

I used the data set provided by the author Long Mingsheng source code.

10 common categories are picked out from 4 datasets of object images: `caltech`, `amazon`, `webcam`, and `dslr`. Then follow the previously reported protocols for preparing features, i.e., **extracting SURF features** and quantizing them into an 800-bin histogram with codebooks computed via K-means on a subset of images from amazon. The file name: `*_SURF_L10.mat` is about features and labels.

I write two version of the JDA.

Refer to Long's

One is the code of *Long Mingsheng's*' MATLAB version of the code. According to the structure of the paper, after the reappearance, the results are as follows:

(**The best accuracy is 79.62%**, about webcam and dslr; but **the worst accuracy is 24.76%**, about webcam and caltech. And except `A->W` and `A->D`, all are better than 1 NN.)

```
root@71b4a3d0c1b1:/home/huangchenxi/test# python JDA.py
[src is ./data/JDA/Caltech10_SURF_L10.mat, tar is ./data/JDA/amazon_SURF_L10.mat
[NN = 0.24217118997912318
=====Iteration [0]=====
[JDA iteration [1/10]: Acc: 0.2777
=====Iteration [1]=====
[JDA iteration [2/10]: Acc: 0.2777
=====Iteration [2]=====
[JDA iteration [3/10]: Acc: 0.2777
=====Iteration [3]=====
[JDA iteration [4/10]: Acc: 0.2777
=====Iteration [4]=====
[JDA iteration [5/10]: Acc: 0.2777
=====Iteration [5]=====
[JDA iteration [6/10]: Acc: 0.2777
=====Iteration [6]=====
[JDA iteration [7/10]: Acc: 0.2777
=====Iteration [7]=====
[JDA iteration [8/10]: Acc: 0.2777
=====Iteration [8]=====
[JDA iteration [9/10]: Acc: 0.2777
=====Iteration [9]=====
[JDA iteration [10/10]: Acc: 0.2777

src is ./data/JDA/Caltech10_SURF_L10.mat, tar is ./data/JDA/webcam_SURF_L10.mat
NN = 0.28135593220338984
=====Iteration [0]=====
[JDA iteration [1/10]: Acc: 0.2847
=====Iteration [1]=====
[JDA iteration [2/10]: Acc: 0.2847
=====Iteration [2]=====
[JDA iteration [3/10]: Acc: 0.2847
=====Iteration [3]=====
[JDA iteration [4/10]: Acc: 0.2847
=====Iteration [4]=====
[JDA iteration [5/10]: Acc: 0.2847
=====Iteration [5]=====
[JDA iteration [6/10]: Acc: 0.2847
=====Iteration [6]=====
[JDA iteration [7/10]: Acc: 0.2847
=====Iteration [7]=====
[JDA iteration [8/10]: Acc: 0.2847
=====Iteration [8]=====
[JDA iteration [9/10]: Acc: 0.2847
=====Iteration [9]=====
[JDA iteration [10/10]: Acc: 0.2847

src is ./data/JDA/Caltech10_SURF_L10.mat, tar is ./data/JDA/dslr_SURF_L10.mat
NN = 0.28662420382165604
=====Iteration [0]=====
[JDA iteration [1/10]: Acc: 0.3694
=====Iteration [1]=====
[JDA iteration [2/10]: Acc: 0.3694
=====Iteration [2]=====
[JDA iteration [3/10]: Acc: 0.3694
=====Iteration [3]=====
[JDA iteration [4/10]: Acc: 0.3694
=====Iteration [4]=====
[JDA iteration [5/10]: Acc: 0.3694
=====Iteration [5]=====
[JDA iteration [6/10]: Acc: 0.3694
=====Iteration [6]=====
[JDA iteration [7/10]: Acc: 0.3694
=====Iteration [7]=====
[JDA iteration [8/10]: Acc: 0.3694
=====Iteration [8]=====
[JDA iteration [9/10]: Acc: 0.3694
=====Iteration [9]=====
[JDA iteration [10/10]: Acc: 0.3694
```

```
src is ./data/JDA/amazon_SURF_L10.mat, tar is ./data/JDA/Caltech10_SURF_L10.mat
NN =  0.15761353517364202
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2680
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2680
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2680
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2680
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2680
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2680
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2680
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2680
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.2680
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2680
```

```
src is ./data/JDA/amazon_SURF_L10.mat, tar is ./data/JDA/webcam_SURF_L10.mat
NN =  0.3152542372881356
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.3085
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.3085
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.3085
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.3085
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.3085
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.3085
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.3085
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.3085
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.3085
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.3085
```

```
src is ./data/JDA/amazon_SURF_L10.mat, tar is ./data/JDA/dslr_SURF_L10.mat
NN =  0.29936305732484075
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2866
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2866
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2866
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2866
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2866
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2866
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2866
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2866
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.2866
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2866
```

```
src is ./data/JDA/webcam_SURF_L10.mat, tar is ./data/JDA/Caltech10_SURF_L10.mat
NN = 0.1772039180765806
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2476
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2476
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2476
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2476
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2476
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2476
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2476
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2476
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.2476
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2476
```

```
src is ./data/JDA/webcam_SURF_L10.mat, tar is ./data/JDA/amazon_SURF_L10.mat
NN = 0.17118997912317327
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2777
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2777
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2777
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2777
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2777
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2777
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2777
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2777
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.2777
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2777
```

```
src is ./data/JDA/webcam_SURF_L10.mat, tar is ./data/JDA/dslr_SURF_L10.mat
NN = 0.6242038216560509
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.7962
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.7962
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.7962
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.7962
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.7962
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.7962
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.7962
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.7962
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.7962
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.7962
```

```
src is ./data/JDA/dslr_SURF_L10.mat, tar is ./data/JDA/Caltech10_SURF_L10.mat
NN = 0.08370436331255565
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2850
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2850
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2850
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2850
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2850
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2850
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2850
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2850
=====Iteration [8]=====
nvidijda iteration [9/10]: Acc: 0.2850
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2850
```

```
src is ./data/JDA/dslr_SURF_L10.mat, tar is ./data/JDA/amazon_SURF_L10.mat
NN = 0.10438413361169102
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.2766
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.2766
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.2766
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.2766
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.2766
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.2766
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.2766
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.2766
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.2766
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.2766
```

```
src is ./data/JDA/dslr_SURF_L10.mat, tar is ./data/JDA/webcam_SURF_L10.mat
NN = 0.10847457627118644
=====Iteration [0]=====
JDA iteration [1/10]: Acc: 0.7356
=====Iteration [1]=====
JDA iteration [2/10]: Acc: 0.7356
=====Iteration [2]=====
JDA iteration [3/10]: Acc: 0.7356
=====Iteration [3]=====
JDA iteration [4/10]: Acc: 0.7356
=====Iteration [4]=====
JDA iteration [5/10]: Acc: 0.7356
=====Iteration [5]=====
JDA iteration [6/10]: Acc: 0.7356
=====Iteration [6]=====
JDA iteration [7/10]: Acc: 0.7356
=====Iteration [7]=====
JDA iteration [8/10]: Acc: 0.7356
=====Iteration [8]=====
JDA iteration [9/10]: Acc: 0.7356
=====Iteration [9]=====
JDA iteration [10/10]: Acc: 0.7356
```

```
root@71b4a3d0c1b1:/home/huangchenxi/test#
```

Compare

First, I think the low accuracy due to my programming, but after I see the paper result, it may because the data or algorithm.

The best in paper is $D \rightarrow W = 89.49\%$, and in my test is $W \rightarrow D = 79.62\%$ (and $D \rightarrow W = 73.85\%$).

The worst in paper is $W \rightarrow C = 31.17\%$, and in my test is $W \rightarrow C = 24.76\%$.

Table 3. Accuracy (%) on 4 types of cross-domain image datasets.

Dataset	Standard Learning		Transfer Learning			
	NN	PCA	GFK	TCA	TSL	JDA
USPS vs MNIST	44.7	44.95	46.45	51.05	53.75	59.65
MNIST vs USPS	65.94	66.22	67.22	56.28	66.06	67.28
COIL1 vs COIL2	83.61	84.72	72.50	88.47	88.06	89.31
COIL2 vs COIL1	82.78	84.03	74.17	85.83	87.92	88.47
PIE1 vs PIE2 (1)	26.09	24.80	26.15	40.76	44.08	58.81
PIE1 vs PIE3 (2)	26.59	25.18	27.27	41.79	47.49	54.23
PIE1 vs PIE4 (3)	30.67	29.26	31.15	59.63	62.78	84.50
PIE1 vs PIE5 (4)	16.67	16.30	17.59	29.35	36.15	49.75
PIE2 vs PIE1 (5)	24.49	24.22	25.24	41.81	46.28	57.62
PIE2 vs PIE3 (6)	46.63	45.53	47.37	51.47	57.60	62.93
PIE2 vs PIE4 (7)	54.07	53.35	54.25	64.73	71.43	75.82
PIE2 vs PIE5 (8)	26.53	25.43	27.08	33.70	35.66	39.89
PIE3 vs PIE1 (9)	21.37	20.95	21.82	34.69	36.94	50.96
PIE3 vs PIE2 (10)	41.01	40.45	43.16	47.70	47.02	57.95
PIE3 vs PIE4 (11)	46.53	46.14	46.41	56.23	59.45	68.45
PIE3 vs PIE5 (12)	26.23	25.31	26.78	33.15	36.34	39.95
PIE4 vs PIE1 (13)	32.95	31.96	34.24	55.64	63.66	80.58
PIE4 vs PIE2 (14)	62.68	60.96	62.92	67.83	72.68	82.63
PIE4 vs PIE3 (15)	73.22	72.18	73.35	75.86	83.52	87.25
PIE4 vs PIE5 (16)	37.19	35.11	37.38	40.26	44.79	54.66
PIE5 vs PIE1 (17)	18.49	18.85	20.35	26.98	33.28	46.46
PIE5 vs PIE2 (18)	24.19	23.39	24.62	29.90	34.13	42.05
PIE5 vs PIE3 (19)	28.31	27.21	28.49	29.90	36.58	53.31
PIE5 vs PIE4 (20)	31.24	30.34	31.33	33.64	38.75	57.01
C → A (1)	23.70	36.95	41.02	38.20	44.47	44.78
C → W (2)	25.76	32.54	40.68	38.64	34.24	41.69
C → D (3)	25.48	38.22	38.85	41.40	43.31	45.22
A → C (4)	26.00	34.73	40.25	37.76	37.58	39.36
A → W (5)	29.83	35.59	38.98	37.63	33.90	37.97
A → D (6)	25.48	27.39	36.31	33.12	26.11	39.49
W → C (7)	19.86	26.36	30.72	29.30	29.83	31.17
W → A (8)	22.96	31.00	29.75	30.06	30.27	32.78
W → D (9)	59.24	77.07	80.89	87.26	87.26	89.17
D → C (10)	26.27	29.65	30.28	31.70	28.50	31.52
D → A (11)	28.50	32.05	32.05	32.15	27.56	33.09
D → W (12)	63.39	75.93	75.59	86.10	85.42	89.49
Average	37.46	39.84	41.19	47.22	49.80	57.37

	In my test(NN)	In paper(NN)	In my test(JDA)	In paper(JDA)
C->A	24.22	23.70	27.77	44.78
C->W	28.14	25.76	28.47	41.69
C->D	28.66	25.48	36.94	45.22
A->C	15.76	26.00	26.80	39.36
A->W	31.53	29.83	30.85	37.97
A->D	29.94	25.48	28.66	39.49
W->C	17.72	19.86	24.76	31.17
W->A	17.12	22.96	27.77	32.78
W->D	62.42	59.24	79.62	89.17
D->C	8.37	26.27	28.50	31.52
D->A	10.44	28.50	27.66	33.09
D->W	10.85	63.39	73.56	89.49

It seems that this reappearance is not good enough, and the gap is a bit big. The first reason is that the data itself is processed and loses some information; Another reason is that matlab is converted to python, and may not use a better function.

Refer to Wang's

The other is *Wang Jindong's* Python version of the code. He changed the structure of the code, reduced the function call, converted to a loop, used the data already obtained in the previous loop, avoiding the recalculation, and add fine-tune. Therefore, the error is also reduced and **is similiar with that in paper**. The results are as follows:

(The accuracy is **46.56%** about caltech and amazon. And in last version is only 27.78%, in paper is 44.78%)

```
[root@71b4a3d0c1b1:/home.huangchenxi# python JDA_2.py
JDA: data = Office-31 lambda = 1 src: data/Office-31/caltech.mat tar: data/Office-31/amazon.mat
JDA iteration [1/10]: Acc: 0.4666
JDA iteration [2/10]: Acc: 0.4593
JDA iteration [3/10]: Acc: 0.4656
JDA iteration [4/10]: Acc: 0.4624
JDA iteration [5/10]: Acc: 0.4666
JDA iteration [6/10]: Acc: 0.4666
JDA iteration [7/10]: Acc: 0.4656
JDA iteration [8/10]: Acc: 0.4656
JDA iteration [9/10]: Acc: 0.4656
JDA iteration [10/10]: Acc: 0.4656
0.46555323590814196
root@71b4a3d0c1b1:/home.huangchenxi#
```

Compare

And the summary in tables:

	In my test	In paper
C->A	46.56	44.78
C->W	39.32	41.69
C->D	47.13	45.22
A->C	40.16	39.36
A->W	38.30	37.97
A->D	36.94	39.49
W->C	32.41	31.17
W->A	32.99	32.78
W->D	92.36	89.17
D->C	32.06	31.52
D->A	33.82	33.09
D->W	89.83	89.49

Compare with TCA

It can be seen that **JDA is basically better than TCA** (it is good in the paper, probably because the degree of fine-tune is different).

	In my test(JDA)	In paper(JDA)	In my test(TCA)	In paper(TCA)
C->A	46.56	44.78	45.62	44.47
C->W	39.32	41.69	39.32	34.24
C->D	47.13	45.22	45.86	43.31
A->C	40.16	39.36	42.03	37.58
A->W	38.30	37.97	40.00	33.90
A->D	36.94	39.49	35.67	26.11
W->C	32.41	31.17	31.52	29.83
W->A	32.99	32.78	30.48	30.27
W->D	92.36	89.17	91.08	87.26
D->C	32.06	31.52	32.95	28.50
D->A	33.82	33.09	32.78	27.56
D->W	89.83	89.49	87.46	85.42

Deep transfer learning

ResNet50

The datasets are the same.

The structure of ResNet50 is as follows:

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112					
conv2_x	56×56	$\left[\begin{array}{l} 3 \times 3, 64 \\ 3 \times 3, 64 \end{array} \right] \times 2$	$\left[\begin{array}{l} 3 \times 3, 64 \\ 3 \times 3, 64 \end{array} \right] \times 1$	$\left[\begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$	$\left[\begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$	$\left[\begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$
conv3_x	28×28	$\left[\begin{array}{l} 3 \times 3, 128 \\ 3 \times 3, 128 \end{array} \right] \times 2$	$\left[\begin{array}{l} 3 \times 3, 128 \\ 3 \times 3, 128 \end{array} \right] \times 1$	$\left[\begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 4$	$\left[\begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 4$	$\left[\begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 8$
conv4_x	14×14	$\left[\begin{array}{l} 3 \times 3, 256 \\ 3 \times 3, 256 \end{array} \right] \times 2$	$\left[\begin{array}{l} 3 \times 3, 256 \\ 3 \times 3, 256 \end{array} \right] \times 1$	$\left[\begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 6$	$\left[\begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 23$	$\left[\begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 36$
conv5_x	7×7	$\left[\begin{array}{l} 3 \times 3, 512 \\ 3 \times 3, 512 \end{array} \right] \times 2$	$\left[\begin{array}{l} 3 \times 3, 512 \\ 3 \times 3, 512 \end{array} \right] \times 1$	$\left[\begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$	$\left[\begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$	$\left[\begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$
	1×1			average pool, 1000-d fc, softmax		
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

Table 1. Architectures for ImageNet. Building blocks are shown in brackets (see also Fig. 5), with the numbers of blocks stacked. Down-sampling is performed by conv3_1, conv4_1, and conv5_1 with a stride of 2.

We can build it layer by layer according to its structure; it can also be used

```
torchvision.models.resnet50(pretrained=False, progress=True, **kwargs). I choose  
the latter, and use the pertained parameters (the parameters are in  
. /codes/ResNet50/resnet50-19c8e357.pth , you can use it).
```

But the images is so large and I first set the `batches_size = 256`. So the docker stop, and then I forced the container size as 8 GB and decreased the batches size to 100.

And the part of the result as follows:

```

|root@a28dd28ff56a:/home/huangchenxi/transfer-learning/ResNet50# python Fine
-tune.py
TCA:  data = Office-31 lambda = 1 src: webcam.mat tar: dslr.mat
3288 822
Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [01/100]---src, loss: 1.926839, acc: 0.5620
Epoch: [01/100]---val, loss: 0.978189, acc: 0.6873
Epoch: [01/100]---tar, loss: 0.990008, acc: 0.6854

Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [02/100]---src, loss: 0.746554, acc: 0.7509
Epoch: [02/100]---val, loss: 0.626346, acc: 0.7032
Epoch: [02/100]---tar, loss: 0.631439, acc: 0.6981

Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [03/100]---src, loss: 0.503822, acc: 0.7859
Epoch: [03/100]---val, loss: 0.416240, acc: 0.8698
Epoch: [03/100]---tar, loss: 0.402277, acc: 0.8832

Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [04/100]---src, loss: 0.364801, acc: 0.8863
Epoch: [04/100]---val, loss: 0.316410, acc: 0.9136
Epoch: [04/100]---tar, loss: 0.302642, acc: 0.9358

Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [05/100]---src, loss: 0.287266, acc: 0.9389
Epoch: [05/100]---val, loss: 0.273787, acc: 0.9294
Epoch: [05/100]---tar, loss: 0.244674, acc: 0.9564

Learning rate: 0.00010000
Learning rate: 0.00100000
Epoch: [06/100]---src, loss: 0.244297, acc: 0.9519
Epoch: [06/100]---val, loss: 0.230979, acc: 0.9574
Epoch: [06/100]---tar, loss: 0.212978, acc: 0.9676

```

Obviously, the accuracy is increase. And only after 4 iteration, the target accuracy can **be up to 93.58%**. So, the network is good. And you can set up the batches size as 10, the loss can even lower.

Compare

Since this part directly calls the function, it does not look at the paper that presented it. So no comparison is made. But with the accuracy of non-deep networks, ResNet has a good performance.

With Wasserstein

Replaced the metrics and replaced MMD with Wasserstein.

Compared with the original form of the original GAN, WGAN only changed four points:

- The last layer of the discriminator removes sigmoid
- The loss of the generator and discriminator does not take the log
- Each time the parameters of the discriminator are updated, their absolute values are

truncated to no more than a fixed constant c

- Do not use momentum-based optimization algorithms (including momentum and Adam), recommend RMSProp, SGD.

The results as follows:

```
[\$ python Baseline.py
Epoch: [01/25]---src, loss: 0.342123, acc: 0.9092
Epoch: [01/25]---val, loss: 0.333888, acc: 0.9200
Epoch: [01/25]---tar, loss: 0.362404, acc: 0.8432

Epoch: [02/25]---src, loss: 0.330124, acc: 0.9036
Epoch: [02/25]---val, loss: 0.333967, acc: 0.9036
Epoch: [02/25]---tar, loss: 0.365839, acc: 0.8653

Epoch: [03/25]---src, loss: 0.330052, acc: 0.9015
Epoch: [03/25]---val, loss: 0.333897, acc: 0.8848
Epoch: [03/25]---tar, loss: 0.364945, acc: 0.8436

Epoch: [04/25]---src, loss: 0.329205, acc: 0.9012
Epoch: [04/25]---val, loss: 0.339296, acc: 0.9222
Epoch: [04/25]---tar, loss: 0.376061, acc: 0.9002

Epoch: [05/25]---src, loss: 0.329224, acc: 0.8982
Epoch: [05/25]---val, loss: 0.333891, acc: 0.8990
Epoch: [05/25]---tar, loss: 0.359992, acc: 0.9098

Epoch: [06/25]---src, loss: 0.328851, acc: 0.8983
Epoch: [06/25]---val, loss: 0.334494, acc: 0.9108
Epoch: [06/25]---tar, loss: 0.357882, acc: 0.9240

Epoch: [07/25]---src, loss: 0.330282, acc: 0.8969
Epoch: [07/25]---val, loss: 0.334287, acc: 0.9284
Epoch: [07/25]---tar, loss: 0.362293, acc: 0.9247

Epoch: [08/25]---src, loss: 0.329142, acc: 0.8970
Epoch: [08/25]---val, loss: 0.345426, acc: 0.9111
Epoch: [08/25]---tar, loss: 0.378538, acc: 0.9237

Epoch: [09/25]---src, loss: 0.328856, acc: 0.9002
Epoch: [09/25]---val, loss: 0.341286, acc: 0.8779
Epoch: [09/25]---tar, loss: 0.362307, acc: 0.7775

Epoch: [10/25]---src, loss: 0.328975, acc: 0.8982
Epoch: [10/25]---val, loss: 0.337882, acc: 0.8985
Epoch: [10/25]---tar, loss: 0.362407, acc: 0.8837

Epoch: [11/25]---src, loss: 0.329133, acc: 0.8991]
```

Relatively speaking, the rate is very fast at first, but after many iterations, the accuracy is little lower(May be the dataset is too small).

Digit Network

After reading Appendix D of "Self-Ensembling for Visual Domain Adaptation". I began to reproduce his digital migration network. I chose the simplest one. That is, the transfer learning from MNIST to USPS is realized. Its network structure is set as follows:

Description	Shape
28 × 28 Mono image	$28 \times 28 \times 1$
Conv $5 \times 5 \times 32$, batch norm	$24 \times 24 \times 32$
Max-pool, 2x2	$12 \times 12 \times 32$
Conv $3 \times 3 \times 64$, batch norm	$10 \times 10 \times 64$
Conv $3 \times 3 \times 64$, batch norm	$8 \times 8 \times 64$
Max-pool, 2x2	$4 \times 4 \times 64$
Dropout, 50%	$4 \times 4 \times 64$
Fully connected, 256 units	256
Fully connected, 10 units, softmax	10

Table 6: MNIST \leftrightarrow USPS architecture

The code of network are as follows:

```
def __init__(self, n_classes):
    self.conv1_1 = nn.Conv2d(1, 32, (5, 5))
    self.conv1_1_bn = nn.BatchNorm2d(32)

    self.pool1 = nn.MaxPool2d((2, 2))

    self.conv2_1 = nn.Conv2d(32, 64, (3, 3))
    self.conv2_1_bn = nn.BatchNorm2d(64)

    self.conv2_2 = nn.Conv2d(64, 64, (3, 3))
    self.conv2_2_bn = nn.BatchNorm2d(64)

    self.pool2 = nn.MaxPool2d((2, 2))

    self.drop1 = nn.Dropout()

    self.fc3 = nn.Linear(1024, 256)

    self.fc4 = nn.Linear(256, n_classes)
```

And the result is:

```
batch size: 256 number workers: 4 learning rate=0.001 epochs: 25
```

```
root@a28dd28ff56a:/home/huangchenxi/transfer-learning/Mnist2Usps# python Baseline.py
Epoch: [01/25]---src, loss: 0.192084, acc: 0.9433
Epoch: [01/25]---val, loss: 0.051905, acc: 0.9825
Epoch: [01/25]---tar, loss: 1.194448, acc: 0.7043

Epoch: [02/25]---src, loss: 0.060604, acc: 0.9806
Epoch: [02/25]---val, loss: 0.046370, acc: 0.9858
Epoch: [02/25]---tar, loss: 1.519606, acc: 0.6295

Epoch: [03/25]---src, loss: 0.047869, acc: 0.9850
Epoch: [03/25]---val, loss: 0.028589, acc: 0.9905
Epoch: [03/25]---tar, loss: 1.143164, acc: 0.7022

Epoch: [04/25]---src, loss: 0.038444, acc: 0.9876
Epoch: [04/25]---val, loss: 0.021601, acc: 0.9928
Epoch: [04/25]---tar, loss: 0.903657, acc: 0.7644

Epoch: [05/25]---src, loss: 0.032450, acc: 0.9901
Epoch: [05/25]---val, loss: 0.039623, acc: 0.9874
Epoch: [05/25]---tar, loss: 1.234952, acc: 0.7335

Epoch: [06/25]---src, loss: 0.028346, acc: 0.9912
Epoch: [06/25]---val, loss: 0.018120, acc: 0.9939
Epoch: [06/25]---tar, loss: 0.971560, acc: 0.7553

Epoch: [07/25]---src, loss: 0.025108, acc: 0.9915
Epoch: [07/25]---val, loss: 0.021287, acc: 0.9933
Epoch: [07/25]---tar, loss: 0.861253, acc: 0.7670

Epoch: [08/25]---src, loss: 0.024066, acc: 0.9922
Epoch: [08/25]---val, loss: 0.021791, acc: 0.9929
Epoch: [08/25]---tar, loss: 1.067999, acc: 0.7469

Epoch: [09/25]---src, loss: 0.020962, acc: 0.9934
Epoch: [09/25]---val, loss: 0.017854, acc: 0.9944
Epoch: [09/25]---tar, loss: 0.802467, acc: 0.7923

Epoch: [10/25]---src, loss: 0.017997, acc: 0.9940
Epoch: [10/25]---val, loss: 0.018616, acc: 0.9936
Epoch: [10/25]---tar, loss: 0.971246, acc: 0.7475

Epoch: [11/25]---src, loss: 0.018466, acc: 0.9939
Epoch: [11/25]---val, loss: 0.018829, acc: 0.9948
Epoch: [11/25]---tar, loss: 0.823219, acc: 0.7963

Epoch: [12/25]---src, loss: 0.016613, acc: 0.9945
Epoch: [12/25]---val, loss: 0.017613, acc: 0.9950
Epoch: [12/25]---tar, loss: 0.801601, acc: 0.8118
```

```
Epoch: [13/25]---src, loss: 0.015181, acc: 0.9949
Epoch: [13/25]---val, loss: 0.018595, acc: 0.9939
Epoch: [13/25]---tar, loss: 0.847482, acc: 0.7943

Epoch: [14/25]---src, loss: 0.014008, acc: 0.9954
Epoch: [14/25]---val, loss: 0.017293, acc: 0.9948
Epoch: [14/25]---tar, loss: 0.938279, acc: 0.7726

Epoch: [15/25]---src, loss: 0.012149, acc: 0.9960
Epoch: [15/25]---val, loss: 0.019157, acc: 0.9940
Epoch: [15/25]---tar, loss: 0.856247, acc: 0.7869

Epoch: [16/25]---src, loss: 0.010967, acc: 0.9964
Epoch: [16/25]---val, loss: 0.020897, acc: 0.9937
Epoch: [16/25]---tar, loss: 1.030946, acc: 0.7727

Epoch: [17/25]---src, loss: 0.011991, acc: 0.9960
Epoch: [17/25]---val, loss: 0.018124, acc: 0.9945
Epoch: [17/25]---tar, loss: 0.753240, acc: 0.8066

Epoch: [18/25]---src, loss: 0.012081, acc: 0.9959
Epoch: [18/25]---val, loss: 0.019345, acc: 0.9941
Epoch: [18/25]---tar, loss: 1.005149, acc: 0.7689

Epoch: [19/25]---src, loss: 0.010824, acc: 0.9962
Epoch: [19/25]---val, loss: 0.016516, acc: 0.9945
Epoch: [19/25]---tar, loss: 0.724475, acc: 0.8185

Epoch: [20/25]---src, loss: 0.009383, acc: 0.9966
Epoch: [20/25]---val, loss: 0.024679, acc: 0.9937
Epoch: [20/25]---tar, loss: 0.885000, acc: 0.7923

Epoch: [21/25]---src, loss: 0.009041, acc: 0.9970
Epoch: [21/25]---val, loss: 0.018507, acc: 0.9948
Epoch: [21/25]---tar, loss: 0.968110, acc: 0.7711

Epoch: [22/25]---src, loss: 0.008792, acc: 0.9969
Epoch: [22/25]---val, loss: 0.020892, acc: 0.9944
Epoch: [22/25]---tar, loss: 0.817327, acc: 0.7844

Epoch: [23/25]---src, loss: 0.009098, acc: 0.9969
Epoch: [23/25]---val, loss: 0.021970, acc: 0.9936
Epoch: [23/25]---tar, loss: 1.027080, acc: 0.7816

Epoch: [24/25]---src, loss: 0.008881, acc: 0.9968
Epoch: [24/25]---val, loss: 0.018496, acc: 0.9945
Epoch: [24/25]---tar, loss: 0.681893, acc: 0.8276

Epoch: [25/25]---src, loss: 0.007519, acc: 0.9976
Epoch: [25/25]---val, loss: 0.019360, acc: 0.9945
Epoch: [25/25]---tar, loss: 0.962274, acc: 0.7816
```

In experience, lowering some learning rates, increasing the number of iterations, improving the size of the batch, etc., can improve some accuracy. The experiment also proved.

```
batch size: 36 number workers: 2 learning rate=3e-4 epochs: 256
```

```
$ python main.py --model='Baseline_Digit' --source='MNIST' --target='USPS' --n_classes=10
Baseline_Digit
Epoch: [01/256]---src, loss: 0.173502, acc: 0.9516
Epoch: [01/256]---val, loss: 0.043602, acc: 0.9847
Epoch: [01/256]---tar, loss: 1.286771, acc: 0.6582

Epoch: [02/256]---src, loss: 0.061344, acc: 0.9817
Epoch: [02/256]---val, loss: 0.030976, acc: 0.9902
Epoch: [02/256]---tar, loss: 1.113913, acc: 0.7166

Epoch: [03/256]---src, loss: 0.046306, acc: 0.9853
Epoch: [03/256]---val, loss: 0.022454, acc: 0.9926
Epoch: [03/256]---tar, loss: 0.853039, acc: 0.7671

Epoch: [04/256]---src, loss: 0.037732, acc: 0.9879
Epoch: [04/256]---val, loss: 0.024954, acc: 0.9917
Epoch: [04/256]---tar, loss: 0.903350, acc: 0.7910

Epoch: [05/256]---src, loss: 0.035376, acc: 0.9887
Epoch: [05/256]---val, loss: 0.021251, acc: 0.9928
Epoch: [05/256]---tar, loss: 1.204596, acc: 0.6985

Epoch: [06/256]---src, loss: 0.028423, acc: 0.9908
Epoch: [06/256]---val, loss: 0.031079, acc: 0.9900
Epoch: [06/256]---tar, loss: 1.043929, acc: 0.7461

Epoch: [07/256]---src, loss: 0.026714, acc: 0.9914
Epoch: [07/256]---val, loss: 0.027179, acc: 0.9911
Epoch: [07/256]---tar, loss: 0.828751, acc: 0.7917

Epoch: [08/256]---src, loss: 0.023795, acc: 0.9925
Epoch: [08/256]---val, loss: 0.018674, acc: 0.9935
Epoch: [08/256]---tar, loss: 0.924846, acc: 0.7823

Epoch: [09/256]---src, loss: 0.020756, acc: 0.9935
Epoch: [09/256]---val, loss: 0.018090, acc: 0.9935
Epoch: [09/256]---tar, loss: 0.860521, acc: 0.7851

Epoch: [10/256]---src, loss: 0.020786, acc: 0.9933
Epoch: [10/256]---val, loss: 0.017851, acc: 0.9944
Epoch: [10/256]---tar, loss: 0.744768, acc: 0.8303

Epoch: [11/256]---src, loss: 0.017645, acc: 0.9943
Epoch: [11/256]---val, loss: 0.017168, acc: 0.9947
```

```
Epoch: [11/256]---src, loss: 0.017645, acc: 0.9943
[Epoch: [11/256]---val, loss: 0.017168, acc: 0.9947
Epoch: [11/256]---tar, loss: 0.693371, acc: 0.8479

Epoch: [12/256]---src, loss: 0.016330, acc: 0.9947
Epoch: [12/256]---val, loss: 0.018545, acc: 0.9943
Epoch: [12/256]---tar, loss: 0.561949, acc: 0.8546

Epoch: [13/256]---src, loss: 0.015491, acc: 0.9950
Epoch: [13/256]---val, loss: 0.017385, acc: 0.9951
Epoch: [13/256]---tar, loss: 0.701490, acc: 0.8351

[Epoch: [14/256]---src, loss: 0.015387, acc: 0.9948
Epoch: [14/256]---val, loss: 0.018047, acc: 0.9940
Epoch: [14/256]---tar, loss: 0.765911, acc: 0.8305

Epoch: [15/256]---src, loss: 0.013603, acc: 0.9954
Epoch: [15/256]---val, loss: 0.016428, acc: 0.9947
Epoch: [15/256]---tar, loss: 0.556106, acc: 0.8631

Epoch: [16/256]---src, loss: 0.011403, acc: 0.9961
Epoch: [16/256]---val, loss: 0.020806, acc: 0.9941
Epoch: [16/256]---tar, loss: 0.918480, acc: 0.8198
[
Epoch: [17/256]---src, loss: 0.011906, acc: 0.9959
Epoch: [17/256]---val, loss: 0.015828, acc: 0.9956
Epoch: [17/256]---tar, loss: 0.726540, acc: 0.8553

Epoch: [18/256]---src, loss: 0.011253, acc: 0.9963
Epoch: [18/256]---val, loss: 0.017504, acc: 0.9947
Epoch: [18/256]---tar, loss: 0.460645, acc: 0.8756

Epoch: [19/256]---src, loss: 0.008843, acc: 0.9970
Epoch: [19/256]---val, loss: 0.022262, acc: 0.9926
[Epoch: [19/256]---tar, loss: 0.633865, acc: 0.8532

Epoch: [20/256]---src, loss: 0.009553, acc: 0.9966
Epoch: [20/256]---val, loss: 0.020464, acc: 0.9942
Epoch: [20/256]---tar, loss: 0.644801, acc: 0.8441

Epoch: [21/256]---src, loss: 0.009262, acc: 0.9967
Epoch: [21/256]---val, loss: 0.018933, acc: 0.9944
Epoch: [21/256]---tar, loss: 0.674756, acc: 0.8512

Epoch: [22/256]---src, loss: 0.008381, acc: 0.9972
[Epoch: [22/256]---val, loss: 0.020346, acc: 0.9939
Epoch: [22/256]---tar, loss: 0.785033, acc: 0.8259

Epoch: [23/256]---src, loss: 0.008080, acc: 0.9972
Epoch: [23/256]---val, loss: 0.017641, acc: 0.9951
Epoch: [23/256]---tar, loss: 0.786440, acc: 0.8244

Epoch: [24/256]---src, loss: 0.008657, acc: 0.9972
Epoch: [24/256]---val, loss: 0.017334, acc: 0.9950
Epoch: [24/256]---tar, loss: 0.732182, acc: 0.8461

Epoch: [25/256]---src, loss: 0.006765, acc: 0.9977
Epoch: [25/256]---val, loss: 0.022630, acc: 0.9949
Epoch: [25/256]---tar, loss: 0.683339, acc: 0.8404
```

```

Epoch: [26/256]---src, loss: 0.008382, acc: 0.9973
Epoch: [26/256]---val, loss: 0.019544, acc: 0.9942
Epoch: [26/256]---tar, loss: 0.679418, acc: 0.8269

```

Compare

In my first 25 iterations, the best is **86.31%**. Since it does not have a separate MT, **MT+CT** is in the range of 88.14 ± -0.34 , then we can imagine that only the result of MT must be correct.

Experimental results of other data sets are compared in the next heading.

	USPS	MNIST	SVHN	MNIST	CIFAR	STL	Syn Digits	Syn Signs
	MNIST	USPS	MNIST	SVHN	STL	CIFAR	SVHN	GTSRB
TRAIN ON SOURCE								
SupSrc*	77.55 ±0.8	82.03 ±1.16	66.5 ±1.93	25.44 ±2.8	72.84 ±0.61	51.88 ±1.44	86.86 ±0.86	96.95 ±0.36
SupSrc+TF	77.53 ±4.63	95.39 ±0.93	68.65 ±1.5	24.86 ±3.29	75.2 ±0.28	59.06 ±1.02	87.45 ±0.65	97.3 ±0.16
SupSrc+TFA	91.97 ±2.15	96.25 ±0.54	71.73 ±5.73	28.69 ±1.59	75.18 ±0.76	59.38 ±0.58	87.16 ±0.85	98.02 ±0.20
Specific aug. ^b	—	—	—	61.99 ±3.9	—	—	—	—
RevGrad ^a [1]	74.01	91.11	73.91	35.67	66.12	56.91	91.09	88.65
DCRN [2]	73.67	91.8	81.97	40.05	66.37	58.65	—	—
G2A [3]	90.8	92.5	84.70	36.4	—	—	—	—
ADDA [4]	90.1	89.4	76.00	—	—	—	—	—
ATT [5]	—	—	86.20	52.8	—	—	93.1	96.2
SBADA-GAN [6]	97.60	95.04	76.14	61.08	—	—	—	—
ADA [7]	—	—	97.6	—	—	—	91.86	97.66
OUR RESULTS								
MT+TF	98.07 ±2.82	98.26 ±0.11	99.18 ±0.12	13.96 ^c ±4.41	80.08 ±0.25	18.3 ±9.03	15.94 ±0.0	98.63 ±0.09
MT+CT*	92.35 ±8.61	88.14 ±0.34	93.33 ±5.88	33.87 ^c ±4.02	77.53 ±0.11	71.65 ±0.67	96.01 ±0.08	98.53 ±0.15
MT+CT+TF	97.28 ±2.74	98.13 ±0.17	98.64 ±0.42	34.15 ^c ±3.56	79.73 ±0.45	74.24 ±0.46	96.51 ±0.08	98.66 ±0.12
MT+CT+TFA	99.54 ±0.04	98.23 ±0.13	99.26 ±0.05	37.49 ^c ±2.44	80.09 ±0.31	69.86 ±1.97	97.11 ±0.04	99.37 ±0.09
Specific aug. ^b	—	—	—	97.0^c ±0.06	—	—	—	—
TRAIN ON TARGET								
SupTgt*	99.53 ±0.02	97.29 ±0.2	99.59 ±0.08	95.7 ±0.13	67.75 ±2.23	88.86 ±0.38	95.62 ±0.2	98.49 ±0.32
SupTgt+TF	99.62 ±0.04	97.65 ±0.17	99.61 ±0.04	96.19 ±0.1	70.98 ±0.79	89.83 ±0.39	96.18 ±0.09	98.64 ±0.09
SupTgt+TFA	99.62 ±0.03	97.83 ±0.17	99.59 ±0.06	96.65 ±0.11	70.03 ±1.13	90.44 ±0.38	96.59 ±0.09	99.22 ±0.22
Specific aug. ^b	—	—	—	97.16 ±0.05	—	—	—	—

[1] Ganin & Lempitsky (2015), [2] Ghifary et al. (2016), [3] Sankaranarayanan et al. (2017), [4] Tzeng et al. (2017), [5] Saito et al. (2017a), [6] Russo et al. (2017), [7] Haeusser et al. (2017)

^a RevGrad results were available in both Ganin & Lempitsky (2015) and Ghifary et al. (2016); we drew results from both papers to obtain results for all of the experiments shown.

^b MNIST → SVHN specific intensity augmentation as described in Section 4.1

^c MNIST → SVHN experiments used class balance loss.

Table 1: Small image benchmark classification accuracy; each result is presented as *mean ± standard deviation*, computed from 5 independent runs. The abbreviations for components of our models are as follows: MT = mean teacher, CT = confidence thresholding, TF = translation and horizontal flip augmentation, TFA = translation, horizontal flip and affine augmentation, * indicates minimal augmentation.

Self-ensemble Visual Domain Adapt Master

First you need to install some packages, such as batchup, skimage.

- When installing batchup, if you use `pip install batchup`, if the version is too high to prevent installation, you can use `pip install batchup==0.1.0`.
- When installing skimage, the command is `pip install scikit-image`; If the problem of timeout still occurs, modify the image source.

The main highlight of the code is the addition of **Confidence thresholding** and **class balance loss**, and the addition of **data enhancements** based on Mean Teacher and Temporal Ensembling. Where `TRAIN clf loss` is the loss of training classification, and `unsup(tgt)` is the loss of the target under unsupervised conditions. You can see that the test results **are close to the supervised way**.

The input information as follows:

```
root at 641aa73707ea in /home/huangchenxi/transfer-learning/Self_ensemble
$ python self_ensemble.py
Loading MNIST...
MNIST:
train: X.shape=(60000, 1, 28, 28), y.shape=(60000, )
val: X.shape=(0, 1, 28, 28), y.shape=(0, )
test: X.shape=(10000, 1, 28, 28), y.shape=(10000, )
MNIST: train: X.min=0.0, X.max=1.0
Loading USPS...
USPS: train:
X.shape=(7291, 1, 28, 28), y.shape=(7291, )
val: X.shape=(0, 1, 28, 28), y.shape=(0, )
test: X.shape=(2007, 1, 28, 28), y.shape=(2007, )
USPS: train: X.min=0.0, X.max=1.0
=====Loaded data=====
Dataset:
SOURCE Train: X.shape=(60000, 1, 28, 28), y.shape=(60000, )
SOURCE Test: X.shape=(10000, 1, 28, 28), y.shape=(10000, )
TARGET Train: X.shape=(7291, 1, 28, 28)
TARGET Test: X.shape=(2007, 1, 28, 28), y.shape=(2007, )
=====Built network=====
```

And the training result as follows:

```
=====Training=====*** Epoch 0 took 13.36s:TRAIN clf loss=0.678767, unsup (tgt) loss=0.000006, conf mask=0.028%, unsup mask=0.028%; SRC TEST ERR=2.910%, TGT TEST student err=17.838%, TGT TEST teacher err=21.574%*** Epoch 1 took 11.91s:TRAIN clf loss=0.213565, unsup (tgt) loss=0.003726, conf mask=16.330%, unsup mask=16.330%; SRC TEST ERR=2.480%, TGT TEST student err=16.791%, TGT TEST teacher err=15.247%*** Epoch 2 took 11.85s:TRAIN clf loss=0.198579, unsup (tgt) loss=0.010548, conf mask=44.331%, unsup mask=44.331%; SRC TEST ERR=3.450%, TGT TEST student err=10.065%, TGT TEST teacher err=10.912%*** Epoch 3 took 11.91s:TRAIN clf loss=0.165477, unsup (tgt) loss=0.014070, conf mask=56.748%, unsup mask=56.748%; SRC TEST ERR=2.000%, TGT TEST student err=8.271%, TGT TEST teacher err=7.773%*** Epoch 4 took 11.73s:TRAIN clf loss=0.141033, unsup (tgt) loss=0.014656, conf mask=65.030%, unsup mask=65.030%; SRC TEST ERR=1.810%, TGT TEST student err=6.527%, TGT TEST teacher err=5.580%*** Epoch 5 took 12.22s:TRAIN clf loss=0.114963, unsup (tgt) loss=0.015976, conf mask=70.810%, unsup mask=70.810%; SRC TEST ERR=1.350%, TGT TEST student err=6.677%, TGT TEST teacher err=4.833%*** Epoch 6 took 12.36s:TRAIN clf loss=0.112600, unsup (tgt) loss=0.016339, conf mask=75.512%, unsup mask=75.512%; SRC TEST ERR=2.650%, TGT TEST student err=5.082%, TGT TEST teacher err=4.584%*** Epoch 7 took 11.68s:TRAIN clf loss=0.103892, unsup (tgt) loss=0.017079, conf mask=77.959%, unsup mask=77.959%; SRC TEST ERR=0.930%, TGT TEST student err=5.381%, TGT TEST teacher err=4.385%*** Epoch 8 took 12.13s:TRAIN clf loss=0.087609, unsup (tgt) loss=0.016692, conf mask=79.148%, unsup mask=79.148%; SRC TEST ERR=1.260%, TGT TEST student err=4.385%, TGT TEST teacher err=4.235%*** Epoch 9 took 12.24s:TRAIN clf loss=0.096784, unsup (tgt) loss=0.016795, conf mask=81.139%, unsup mask=81.139%; SRC TEST ERR=1.070%, TGT TEST student err=5.531%, TGT TEST teacher err=3.837%*** Epoch 10 took 12.13s:TRAIN clf loss=0.103647, unsup (tgt) loss=0.017304, conf mask=82.826%, unsup mask=82.826%; SRC TEST ERR=1.280%, TGT TEST student err=4.634%, TGT TEST teacher err=3.538%Epoch 11 took 12.16s:TRAIN clf loss=0.080447, unsup (tgt) loss=0.016555, conf mask=81.817%, unsup mask=81.817%; SRC TEST ERR=1.210%, TGT TEST student err=4.634%, TGT TEST teacher err=3.538%Epoch 12 took 12.07s:TRAIN clf loss=0.078208, unsup (tgt) loss=0.016775, conf mask=82.647%, unsup mask=82.647%; SRC TEST ERR=2.050%, TGT TEST student err=4.136%, TGT TEST teacher err=3.438%*** Epoch 13 took 12.00s:TRAIN clf loss=0.079053, unsup (tgt) loss=0.017073, conf mask=84.403%, unsup mask=84.403%; SRC TEST ERR=0.870%, TGT TEST student err=3.737%, TGT TEST teacher err=3.438%*** Epoch 14 took 11.93s:TRAIN clf loss=0.071268, unsup (tgt) loss=0.017110, conf mask=85.246%, unsup mask=85.246%; SRC TEST ERR=1.150%, TGT TEST student err=4.086%, TGT TEST teacher err=2.840%
```

And so on. The numbers are fluctuating, but the overall results are good

```
TRAIN clf loss=0.024455, unsup (tgt) loss=0.017092, conf mask=96.184%, unsup mask=96.184%; SRC TEST ERR=0.640%,
TGT TEST student err=2.242%, TGT TEST teacher err=1.794%
Epoch 185 took 12.51s:
TRAIN clf loss=0.023293, unsup (tgt) loss=0.017368, conf mask=96.114%, unsup mask=96.114%; SRC TEST ERR=0.540%,
TGT TEST student err=1.844%, TGT TEST teacher err=1.794%
*** Epoch 186 took 12.30s:
TRAIN clf loss=0.024465, unsup (tgt) loss=0.017586, conf mask=96.695%, unsup mask=96.695%; SRC TEST ERR=0.600%,
TGT TEST student err=1.943%, TGT TEST teacher err=1.744%
Epoch 187 took 12.53s:
TRAIN clf loss=0.026973, unsup (tgt) loss=0.017184, conf mask=95.534%, unsup mask=95.534%; SRC TEST ERR=0.660%,
TGT TEST student err=1.993%, TGT TEST teacher err=1.744%
Epoch 188 took 12.43s:
TRAIN clf loss=0.018069, unsup (tgt) loss=0.017135, conf mask=95.700%, unsup mask=95.700%; SRC TEST ERR=0.630%,
TGT TEST student err=1.993%, TGT TEST teacher err=1.644%
Epoch 189 took 12.51s:
TRAIN clf loss=0.018568, unsup (tgt) loss=0.017056, conf mask=95.755%, unsup mask=95.755%; SRC TEST ERR=0.490%,
TGT TEST student err=2.093%, TGT TEST teacher err=1.794%
Epoch 190 took 12.44s:
TRAIN clf loss=0.030494, unsup (tgt) loss=0.017192, conf mask=96.128%, unsup mask=96.128%; SRC TEST ERR=0.560%,
TGT TEST student err=1.993%, TGT TEST teacher err=1.694%
Epoch 191 took 12.43s:
TRAIN clf loss=0.015902, unsup (tgt) loss=0.017317, conf mask=96.142%, unsup mask=96.142%; SRC TEST ERR=0.460%,
TGT TEST student err=1.943%, TGT TEST teacher err=1.545%
Epoch 192 took 12.61s:
TRAIN clf loss=0.024231, unsup (tgt) loss=0.017064, conf mask=96.253%, unsup mask=96.253%; SRC TEST ERR=0.520%,
TGT TEST student err=1.993%, TGT TEST teacher err=1.594%
Epoch 193 took 12.49s:
TRAIN clf loss=0.030299, unsup (tgt) loss=0.017117, conf mask=95.783%, unsup mask=95.783%; SRC TEST ERR=0.520%,
TGT TEST student err=1.744%, TGT TEST teacher err=1.545%
Epoch 194 took 12.63s:
TRAIN clf loss=0.017378, unsup (tgt) loss=0.017203, conf mask=95.437%, unsup mask=95.437%; SRC TEST ERR=0.440%,
TGT TEST student err=1.893%, TGT TEST teacher err=1.594%
Epoch 195 took 12.44s:
TRAIN clf loss=0.018870, unsup (tgt) loss=0.017045, conf mask=96.142%, unsup mask=96.142%; SRC TEST ERR=0.680%,
TGT TEST student err=1.943%, TGT TEST teacher err=1.545%
Epoch 196 took 12.35s:
TRAIN clf loss=0.022651, unsup (tgt) loss=0.017567, conf mask=96.225%, unsup mask=96.225%; SRC TEST ERR=0.530%,
TGT TEST student err=1.844%, TGT TEST teacher err=1.744%
Epoch 197 took 12.72s:
TRAIN clf loss=0.023108, unsup (tgt) loss=0.017201, conf mask=96.336%, unsup mask=96.336%; SRC TEST ERR=0.590%,
TGT TEST student err=2.043%, TGT TEST teacher err=1.844%
Epoch 198 took 12.31s:
TRAIN clf loss=0.022707, unsup (tgt) loss=0.017238, conf mask=95.893%, unsup mask=95.893%; SRC TEST ERR=0.530%,
TGT TEST student err=1.893%, TGT TEST teacher err=1.594%
Epoch 199 took 12.51s:
TRAIN clf loss=0.016236, unsup (tgt) loss=0.017005, conf mask=95.796%, unsup mask=95.796%; SRC TEST ERR=0.680%,
TGT TEST student err=1.993%, TGT TEST teacher err=1.744%
Closing remaining open files:/root/.batchup/data/mnist.h5...done/root/.batchup/data/usps.h5...done
```

After 200 epochs, the result is respectively perfect.

Obviously, the loss is smaller, and you can see the target loss in unsupervised is so small that can even **be same with supervised**.

Compare

The results of the paper is:

	USPS	MNIST	SVHN	MNIST	CIFAR	STL	Syn Digits	Syn Signs
	MNIST	USPS	MNIST	SVHN	STL	CIFAR	SVHN	GTSRB
TRAIN ON SOURCE								
SupSrc [*]	77.55 ±0.8	82.03 ±1.16	66.5 ±1.93	25.44 ±2.8	72.84 ±0.61	51.88 ±1.44	86.86 ±0.86	96.95 ±0.36
SupSrc+TF	77.53 ±4.63	95.39 ±0.93	68.65 ±1.5	24.86 ±3.29	75.2 ±0.28	59.06 ±1.02	87.45 ±0.65	97.3 ±0.16
SupSrc+TFA	91.97 ±2.15	96.25 ±0.54	71.73 ±5.73	28.69 ±1.59	75.18 ±0.76	59.38 ±0.58	87.16 ±0.85	98.02 ±0.20
Specific aug. ^b	—	—	—	61.99 ±3.9	—	—	—	—
RevGrad ^a [1]	74.01	91.11	73.91	35.67	66.12	56.91	91.09	88.65
DCRN [2]	73.67	91.8	81.97	40.05	66.37	58.65	—	—
G2A [3]	90.8	92.5	84.70	36.4	—	—	—	—
ADDA [4]	90.1	89.4	76.00	—	—	—	—	—
ATT [5]	—	—	86.20	52.8	—	—	93.1	96.2
SBADA-GAN [6]	97.60	95.04	76.14	61.08	—	—	—	—
ADA [7]	—	—	97.6	—	—	—	91.86	97.66
OUR RESULTS								
MT+TF	98.07 ±2.82	98.26 ±0.11	99.18 ±0.12	13.96 ^c ±4.41	80.08 ±0.25	18.3 ±9.03	15.94 ±0.0	98.63 ±0.09
MT+CT [*]	92.35 ±8.61	88.14 ±0.34	93.33 ±5.88	33.87 ^c ±4.02	77.53 ±0.11	71.65 ±0.67	96.01 ±0.08	98.53 ±0.15
MT+CT+TF	97.28 ±2.74	98.13 ±0.17	98.64 ±0.42	34.15 ^c ±3.56	79.73 ±0.45	74.24 ±0.46	96.51 ±0.08	98.66 ±0.12
MT+CT+TFA	99.54 ±0.04	98.23 ±0.13	99.26 ±0.05	37.49 ^c ±2.44	80.09 ±0.31	69.86 ±1.97	97.11 ±0.04	99.37 ±0.09
Specific aug. ^b	—	—	—	97.0^c ±0.06	—	—	—	—
TRAIN ON TARGET								
SupTgt [*]	99.53 ±0.02	97.29 ±0.2	99.59 ±0.08	95.7 ±0.13	67.75 ±2.23	88.86 ±0.38	95.62 ±0.2	98.49 ±0.32
SupTgt+TF	99.62 ±0.04	97.65 ±0.17	99.61 ±0.04	96.19 ±0.1	70.98 ±0.79	89.83 ±0.39	96.18 ±0.09	98.64 ±0.09
SupTgt+TFA	99.62 ±0.03	97.83 ±0.17	99.59 ±0.06	96.65 ±0.11	70.03 ±1.13	90.44 ±0.38	96.59 ±0.09	99.22 ±0.22
Specific aug. ^b	—	—	—	97.16 ±0.05	—	—	—	—

[1] Ganin & Lempitsky (2015), [2] Ghifary et al. (2016), [3] Sankaranarayanan et al. (2017), [4] Tzeng et al. (2017), [5] Saito et al. (2017a), [6] Russo et al. (2017), [7] Haeusser et al. (2017)

^a RevGrad results were available in both Ganin & Lempitsky (2015) and Ghifary et al. (2016); we drew results from both papers to obtain results for all of the experiments shown.

^b MNIST → SVHN specific intensity augmentation as described in Section 4.1.

^c MNIST → SVHN experiments used class balance loss.

Table 1: Small image benchmark classification accuracy; each result is presented as *mean ± standard deviation*, computed from 5 independent runs. The abbreviations for components of our models are as follows: MT = mean teacher, CT = confidence thresholding, TF = translation and horizontal flip augmentation, TFA = translation, horizontal flip and affine augmentation, * indicates minimal augmentation.

I can see that **my error rate is about 1.9%**, that is, **the accuracy rate is about 98%**, which is the same as in the paper. (But my runtime is so long. Maybe something bad I ignore.)

	In paper	In my test
U->M	92.35 ± 8.61	89.75
M->U	88.14 ± 0.34	95.79
S->M	95.33 ± 5.88	91.97
M->S	33.87 ± 4.02	63.23

Other dataset combination tests is in `results_image/Self-ensemble-X-Y.png`. The kinds are follows:

```

if exp == 'svhn_mnist':
    d_source = Data_transform.load_svhn(zero_centre=False, greyscale=True)
    d_target = Data_transform.load_mnist(invert=False, zero_centre=False, pad32=True, val=False)
elif exp == 'mnist_svhn':
    d_source = Data_transform.load_mnist(invert=False, zero_centre=False, pad32=True)
    d_target = Data_transform.load_svhn(zero_centre=False, greyscale=True, val=False)
elif exp == 'svhn_mnist_rgb':
    d_source = Data_transform.load_svhn(zero_centre=False, greyscale=False)
    d_target = Data_transform.load_mnist(invert=False, zero_centre=False, pad32=True, val=False, rgb=True)
elif exp == 'mnist_svhn_rgb':
    d_source = Data_transform.load_mnist(invert=False, zero_centre=False, pad32=True, rgb=True)
    d_target = Data_transform.load_svhn(zero_centre=False, greyscale=False, val=False)
elif exp == 'cifar_stl':
    d_source = Data_transform.load_cifar10(range_01=False)
    d_target = Data_transform.load_stl(zero_centre=False, val=False)
elif exp == 'stl_cifar':
    d_source = Data_transform.load_stl(zero_centre=False)
    d_target = Data_transform.load_cifar10(range_01=False, val=False)
elif exp == 'mnist_usps':
    d_source = Data_transform.load_mnist(zero_centre=False)
    d_target = Data_transform.load_usps(zero_centre=False, scale28=True, val=False)
elif exp == 'usps_mnist':
    d_source = Data_transform.load_usps(zero_centre=False, scale28=True)
    d_target = Data_transform.load_mnist(zero_centre=False, val=False)
elif exp == 'syndigits_svhn':
    d_source = Data_transform.load_syn_digits(zero_centre=False)
    d_target = Data_transform.load_svhn(zero_centre=False, val=False)
elif exp == 'synsigns_gtsrb':
    d_source = Data_transform.load_syn_signs(zero_centre=False)
    d_target = Data_transform.load_gtsrb(zero_centre=False, val=False)

```

Domain Adversarial Training of Neural Networks

The data set should first convert and preprocessed by the [SVMlight toolkit](#). According to the algorithm of the shallow NN, I did [5.1.5 Proxy Distance-Figure 3](#), tested the four combinations:

- whether there is with adversarial
- whether there is with mSDA

The Shallow NN is:

Algorithm: Shallow DANN – Stochastic training update

Input:

1. Samples

$$S = (x_i, y_i)_{i=1}^n, T = (x_i)_{i=1}^{n'}$$

2. hidden layer size D

3. input layer parameter λ

4. learning rate μ

Input: CNN{ W, V, b, c }

$$W, V \leftarrow \text{random_init}(D)$$

$$b, c, u, d \leftarrow 0$$

while stopping criterion is not met do

for i from 1 to n do

// forward propagation

$$G_f(x_i) \leftarrow \text{sigm}(b + Wx_i)$$

$$G_y(G_f(x_i)) \leftarrow \text{softmax}(VG_f(x_i) + c)$$

// backward propagation

$$\Delta_c \leftarrow (e(y_i) - G_y(G_f(x_i)))$$

$$\Delta_V \leftarrow \Delta_c G_f(x_i)^T$$

$$\Delta_b \leftarrow (V^T \Delta_c) \odot G_f(x_i) \odot (1 - G_f(x_i))$$

$$\Delta W \leftarrow \Delta_b \cdot (x_i)^T$$

//from now domain, regularization

$$G_d(G_f(x_i)) \leftarrow \text{sigm}(d + u^T G_f(x_i))$$

$$\Delta_d \leftarrow \lambda(1 - G_d(G_f(x_i)))$$

$$\Delta_u \leftarrow \lambda(1 - G_d(G_f(x_i))) G_f(x_i)$$

$$\text{tmp} \leftarrow \lambda(1 - G_d(G_f(x_i))) \times u \odot G_f(x_i) \odot (1 - G_f(x_i))$$

$$\Delta_b \leftarrow \Delta_b + \text{tmp}$$

$$\Delta_W \leftarrow \Delta_W + \text{tmp} \cdot (x_i)^T$$

//from other domain, regularization

$$j \leftarrow \text{uniform_integer}(1, \dots, n')$$

$$G_f(x_j) \leftarrow \text{sigm}(b + Wx_j)$$

$$G_d(G_f(x_j)) \leftarrow \text{sigm}(d + u^T G_f(x_j))$$

$$\Delta_d \leftarrow \Delta_d - \lambda G_d(G_f(x_j))$$

$$\Delta_u \leftarrow \Delta_u - \lambda G_d(G_f(x_j)) G_f(x_j)$$

$$\text{tmp} \leftarrow -\lambda G_d(G_f(x_j)) \times u \odot G_f(x_j) \odot (1 - G_f(x_j))$$

```

 $\Delta_b \leftarrow \Delta_b + tmp$ 
 $\Delta_W \leftarrow \Delta_W + tmp \cdot (x_j)^T$ 
//update CNN inner parameters
 $W \leftarrow W - \mu \Delta_W$ 
 $V \leftarrow V - \mu \Delta_V$ 
 $b \leftarrow b - \mu \Delta_b$ 
 $c \leftarrow c - \mu \Delta_c$ 
//update domain classifier
 $u \leftarrow u - \mu \Delta_u$ 
 $d \leftarrow d - \mu \Delta_d$ 
end for
end while

```

And the code for the shallow NN is in `DANN.DANN.fit()`.

The data set is `Office-31-Amazon`, the source domain is `DVD`, the target The domain is `electronics`. At the same time, each group has a set of comparisons, whether the `PAD-agent distance method` is used. Three types of training loss, verification loss, and test loss were calculated.

The results are as follows:

- With mSDA
 - Without adversarial

```
Loading data...
source file: ../data/Office-31/Amazon/dvd_train.svmlight
target file: ../data/Office-31/Amazon/electronics_train.svmlight
test file:   ../data/Office-31/Amazon/electronics_test.svmlight
Computing mSDA representations...
Fit...
[DANN parameters] {'hidden_layer_size': 50, 'maxiter': 200, 'lambda_adapt': 0.0, 'epsilon_init': None, 'learning_rate': 0.0001, 'adversarial_representation': False, 'seed': 12342, 'verbose': True}
[DANN best valid risk so far] 0.255000 (iter 0)
[DANN best valid risk so far] 0.240000 (iter 1)
[DANN best valid risk so far] 0.230000 (iter 2)
[DANN best valid risk so far] 0.215000 (iter 3)
[DANN best valid risk so far] 0.210000 (iter 4)
[DANN best valid risk so far] 0.210000 (iter 5)
[DANN best valid risk so far] 0.210000 (iter 6)
[DANN best valid risk so far] 0.210000 (iter 7)
[DANN best valid risk so far] 0.210000 (iter 8)
[DANN best valid risk so far] 0.210000 (iter 9)
[DANN best valid risk so far] 0.210000 (iter 10)
[DANN best valid risk so far] 0.210000 (iter 12)
[DANN best valid risk so far] 0.210000 (iter 13)
[DANN best valid risk so far] 0.210000 (iter 14)
[DANN best valid risk so far] 0.210000 (iter 15)
[DANN best valid risk so far] 0.210000 (iter 16)
[DANN best valid risk so far] 0.210000 (iter 19)
[DANN early stop] iter 31
Predict...
Training Risk    = 0.124444
Validation Risk = 0.210000
Test Risk        = 0.231121
=====
```

```
=====
Computing PAD on DANN representation...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.001000 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.010000 ] train risk: 0.337368 test risk: 0.351053
[ PAD C = 0.100000 ] train risk: 0.146842 test risk: 0.157895
[ PAD C = 1.000000 ] train risk: 0.115789 test risk: 0.122632
[ PAD C = 10.000000 ] train risk: 0.110526 test risk: 0.123684
[ PAD C = 100.000000 ] train risk: 0.112632 test risk: 0.127368
[ PAD C = 1000.000000 ] train risk: 0.108947 test risk: 0.126842
[ PAD C = 10000.000000 ] train risk: 0.108947 test risk: 0.126316
PAD on DANN representation = 1.509474
=====

Computing PAD on original data...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.051053 test risk: 0.060526
[ PAD C = 0.000100 ] train risk: 0.017895 test risk: 0.026842
[ PAD C = 0.001000 ] train risk: 0.012632 test risk: 0.018421
[ PAD C = 0.010000 ] train risk: 0.001053 test risk: 0.021579
[ PAD C = 0.100000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 1.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 10.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 100.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 1000.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 10000.000000 ] train risk: 0.000000 test risk: 0.027895
PAD on original data = 1.926316
```

- With adversarial

```

Loading data...
source file: ../data/Office-31/Amazon/dvd_train.svmlight
target file: ../data/Office-31/Amazon/electronics_train.svmlight
test file:   ../data/Office-31/Amazon/electronics_test.svmlight
mSDA representations loaded from disk
Fit...
[DANN parameters] {'hidden_layer_size': 50, 'maxiter': 200, 'lambda_adapt': 0.1, 'epsilon_init': None, 'learning_rate': 0.0001, 'adversarial_representation': True, 'seed': 12342, 'verbose': True}
[DANN best valid risk so far] 0.255000 (iter 0)
[DANN best valid risk so far] 0.240000 (iter 1)
[DANN best valid risk so far] 0.230000 (iter 2)
[DANN best valid risk so far] 0.215000 (iter 3)
[DANN best valid risk so far] 0.210000 (iter 4)
[DANN best valid risk so far] 0.210000 (iter 5)
[DANN best valid risk so far] 0.210000 (iter 6)
[DANN best valid risk so far] 0.210000 (iter 7)
[DANN best valid risk so far] 0.210000 (iter 8)
[DANN best valid risk so far] 0.210000 (iter 9)
[DANN best valid risk so far] 0.210000 (iter 10)
[DANN best valid risk so far] 0.210000 (iter 12)
[DANN best valid risk so far] 0.210000 (iter 13)
[DANN best valid risk so far] 0.210000 (iter 14)
[DANN best valid risk so far] 0.210000 (iter 15)
[DANN best valid risk so far] 0.210000 (iter 16)
[DANN best valid risk so far] 0.210000 (iter 19)
[DANN early stop] iter 31
Predict...
Training Risk = 0.125000
Validation Risk = 0.210000
Test Risk = 0.223552
=====

```

```

Computing PAD on DANN representation...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.001000 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.010000 ] train risk: 0.362632 test risk: 0.365789
[ PAD C = 0.100000 ] train risk: 0.207368 test risk: 0.221579
[ PAD C = 1.000000 ] train risk: 0.153684 test risk: 0.162105
[ PAD C = 10.000000 ] train risk: 0.144211 test risk: 0.157895
[ PAD C = 100.000000 ] train risk: 0.138947 test risk: 0.154737
[ PAD C = 1000.000000 ] train risk: 0.141579 test risk: 0.154737
[ PAD C = 10000.000000 ] train risk: 0.143158 test risk: 0.157895
PAD on DANN representation = 1.381053
=====
Computing PAD on original data...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.051053 test risk: 0.060526
[ PAD C = 0.000100 ] train risk: 0.017895 test risk: 0.026842
[ PAD C = 0.001000 ] train risk: 0.012632 test risk: 0.018421
[ PAD C = 0.010000 ] train risk: 0.001053 test risk: 0.021579
[ PAD C = 0.100000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 1.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 10.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 100.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 1000.000000 ] train risk: 0.000000 test risk: 0.027895
[ PAD C = 10000.000000 ] train risk: 0.000000 test risk: 0.027895
PAD on original data = 1.926316

```

- Without mSDA

- o Without adversarial

```
>Loading data...
source file: ../data/Office-31/Amazon/dvd_train.svmlight
target file: ../data/Office-31/Amazon/electronics_train.svmlight
test file:   ../data/Office-31/Amazon/electronics_test.svmlight
Fit...
[DANN parameters] {'hidden_layer_size': 50, 'maxiter': 200, 'lambda_adapt': 0.0, 'epsilon_init': None, 'learning_rate': 0.001, 'adversarial_representation': False, 'seed': 12342, 'verbose': True}
[DANN best valid risk so far] 0.365000 (iter 0)
[DANN best valid risk so far] 0.320000 (iter 1)
[DANN best valid risk so far] 0.275000 (iter 2)
[DANN best valid risk so far] 0.260000 (iter 3)
[DANN best valid risk so far] 0.240000 (iter 4)
[DANN best valid risk so far] 0.235000 (iter 5)
[DANN best valid risk so far] 0.220000 (iter 6)
[DANN best valid risk so far] 0.210000 (iter 7)
[DANN best valid risk so far] 0.210000 (iter 8)
[DANN best valid risk so far] 0.195000 (iter 9)
[DANN best valid risk so far] 0.195000 (iter 10)
[DANN best valid risk so far] 0.195000 (iter 12)
[DANN best valid risk so far] 0.195000 (iter 13)
[DANN best valid risk so far] 0.195000 (iter 16)
[DANN best valid risk so far] 0.195000 (iter 17)
[DANN best valid risk so far] 0.195000 (iter 18)
[DANN best valid risk so far] 0.195000 (iter 19)
[DANN best valid risk so far] 0.195000 (iter 21)
[DANN best valid risk so far] 0.195000 (iter 22)
[DANN best valid risk so far] 0.195000 (iter 23)
[DANN best valid risk so far] 0.195000 (iter 24)
[DANN best valid risk so far] 0.195000 (iter 25)
[DANN best valid risk so far] 0.195000 (iter 26)
[DANN best valid risk so far] 0.195000 (iter 27)
[DANN best valid risk so far] 0.195000 (iter 28)
[DANN best valid risk so far] 0.195000 (iter 43)
[DANN best valid risk so far] 0.195000 (iter 44)
[DANN best valid risk so far] 0.195000 (iter 45)
[DANN best valid risk so far] 0.195000 (iter 46)
[DANN best valid risk so far] 0.195000 (iter 70)
[DANN best valid risk so far] 0.195000 (iter 71)
[DANN best valid risk so far] 0.195000 (iter 72)
[DANN best valid risk so far] 0.195000 (iter 87)
```

```

[DANN best valid risk so far] 0.190000 (iter 104)
[DANN best valid risk so far] 0.190000 (iter 151)
[DANN best valid risk so far] 0.190000 (iter 152)
[DANN best valid risk so far] 0.190000 (iter 153)
[DANN best valid risk so far] 0.190000 (iter 154)
[DANN best valid risk so far] 0.190000 (iter 155)
[DANN best valid risk so far] 0.190000 (iter 156)
[DANN best valid risk so far] 0.190000 (iter 157)
[DANN best valid risk so far] 0.190000 (iter 158)
Predict...
Training Risk = 0.000000
Validation Risk = 0.190000
Test Risk = 0.265798
=====
Computing PAD on DANN representation...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.001000 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.010000 ] train risk: 0.471053 test risk: 0.467368
[ PAD C = 0.100000 ] train risk: 0.337368 test risk: 0.337368
[ PAD C = 1.000000 ] train risk: 0.274737 test risk: 0.287895
[ PAD C = 10.000000 ] train risk: 0.268947 test risk: 0.285789
[ PAD C = 100.000000 ] train risk: 0.273684 test risk: 0.287368
[ PAD C = 1000.000000 ] train risk: 0.272632 test risk: 0.288421
[ PAD C = 10000.000000 ] train risk: 0.273684 test risk: 0.287368
PAD on DANN representation = 0.856842
=====
Computing PAD on original data...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.366316 test risk: 0.350526
[ PAD C = 0.001000 ] train risk: 0.104737 test risk: 0.125789
[ PAD C = 0.010000 ] train risk: 0.024737 test risk: 0.052105
[ PAD C = 0.100000 ] train risk: 0.000526 test risk: 0.036842
[ PAD C = 1.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 10.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 100.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 1000.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 10000.000000 ] train risk: 0.000000 test risk: 0.041053
PAD on original data = 1.852632

```

- With adversarial

```
Loading data...
source file: ../data/Office-31/Amazon/dvd_train.svmlight
target file: ../data/Office-31/Amazon/electronics_train.svmlight
test file:   ../data/Office-31/Amazon/electronics_test.svmlight
Fit...
[DANN parameters] {'hidden_layer_size': 50, 'maxiter': 200, 'lambda_adapt': 0.1, 'epsilon_init': None, 'learning_rate': 0.001, 'adversarial_representation': True, 'seed': 12342, 'verbose': True}
[DANN best valid risk so far] 0.365000 (iter 0)
[DANN best valid risk so far] 0.320000 (iter 1) ]
[DANN best valid risk so far] 0.275000 (iter 2)
[DANN best valid risk so far] 0.260000 (iter 3)
[DANN best valid risk so far] 0.240000 (iter 4)
[DANN best valid risk so far] 0.235000 (iter 5)
[DANN best valid risk so far] 0.220000 (iter 6)
[DANN best valid risk so far] 0.210000 (iter 7)
[DANN best valid risk so far] 0.210000 (iter 8)
[DANN best valid risk so far] 0.205000 (iter 9)
[DANN best valid risk so far] 0.195000 (iter 10)
[DANN best valid risk so far] 0.195000 (iter 12)
[DANN best valid risk so far] 0.195000 (iter 13)
[DANN best valid risk so far] 0.195000 (iter 15)
[DANN best valid risk so far] 0.195000 (iter 16)
[DANN best valid risk so far] 0.195000 (iter 17)
[DANN best valid risk so far] 0.195000 (iter 23) ]
[DANN best valid risk so far] 0.195000 (iter 24)
[DANN best valid risk so far] 0.195000 (iter 25)
[DANN best valid risk so far] 0.195000 (iter 26)
[DANN best valid risk so far] 0.195000 (iter 27)
[DANN best valid risk so far] 0.195000 (iter 28)
[DANN best valid risk so far] 0.195000 (iter 29)
[DANN best valid risk so far] 0.195000 (iter 42)
[DANN early stop] iter 64
Predict...
Training Risk = 0.026667
Validation Risk = 0.195000
Test Risk      = 0.254357
=====
```

```
=====
Computing PAD on DANN representation...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.001000 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.010000 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.100000 ] train risk: 0.372105 test risk: 0.379474
[ PAD C = 1.000000 ] train risk: 0.265263 test risk: 0.285263
[ PAD C = 10.000000 ] train risk: 0.265263 test risk: 0.279474
[ PAD C = 100.000000 ] train risk: 0.265789 test risk: 0.277895
[ PAD C = 1000.000000 ] train risk: 0.264737 test risk: 0.276842
[ PAD C = 10000.000000 ] train risk: 0.263158 test risk: 0.277368
PAD on DANN representation = 0.892632
=====

Computing PAD on original data...
PAD on (1800, 2000) examples
[ PAD C = 0.000010 ] train risk: 0.473684 test risk: 0.473684
[ PAD C = 0.000100 ] train risk: 0.366316 test risk: 0.350526
[ PAD C = 0.001000 ] train risk: 0.104737 test risk: 0.125789
[ PAD C = 0.010000 ] train risk: 0.024737 test risk: 0.052105
[ PAD C = 0.100000 ] train risk: 0.000526 test risk: 0.036842
[ PAD C = 1.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 10.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 100.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 1000.000000 ] train risk: 0.000000 test risk: 0.041053
[ PAD C = 10000.000000 ] train risk: 0.000000 test risk: 0.041053
PAD on original data = 1.852632
```

Compare

Let's look at the results in tabular form for comparison.

		Without adversarial	With adversarial
With mSDA	Training Risk	0.124444	0.125000
	Validation Risk	0.210000	0.210000
	Test Risk	0.231121	0.223552
	PAD on DANN	1.509474	1.381053
	PAD on original data	1.926316	1.926316
	Iteration numbers	19	19
Without mSDA	Training Risk	0.000000	0.026667
	Validation Risk	0.190000	0.195000
	Test Risk	0.265798	0.254357
	PAD on DANN	0.856842	0.892632
	PAD on original data	1.852632	1.852632
	Iteration numbers	158	42

In paper: D->E

DANN on Original data: ≈ 1.25

DANN & NN with 100 hidden neurons: ≈ 0.6

DANN on mSDA representation: ≈ 1.0

First we can see that mSDA can accelerate the process of convergence and also can increase PAD. (PAD is a metric estimating the **similarity** of the source and the target representations.)

In the absence of mSDA and no adversarial, although there is no error in the training set, in the test set, the error is relatively large and there is a tendency to overfitting.

In these four cases, **when there is both mSDA and confrontation**, the effect is best. The convergence speed is fast and the test results are good.

But at the same time, we can also see that there is no difference of more than 0.02 between the test results in the four cases. Therefore, it also illustrates **the generalization of DANN**. And its design is simple and can be attached to many previous models to improve performance.

ADA

Most important is the `associated loss`.

```
class WalkerLoss(nn.Module):

    def forward(self, Psts, y):
        equality_matrix = torch.eq(y.clone().view(-1, 1), y).float()
        p_target = equality_matrix / equality_matrix.sum(dim=1, keepdim=True)
        p_target.requires_grad = False

        L_walker = F.kl_div(torch.log(1e-8 + Psts), p_target,
size_average=False)
        L_walker /= p_target.size()[0]

    return L_walker


class VisitLoss(nn.Module):

    def forward(self, Pt):
        p_visit = torch.ones([1, Pt.size()[1]]) / float(Pt.size()[1])
        p_visit.requires_grad = False
        if Pt.is_cuda: p_visit = p_visit.cuda()
        L_visit = F.kl_div(torch.log(1e-8 + Pt), p_visit, size_average=False)
        L_visit /= p_visit.size()[0]

    return L_visit
```

Due to the space problem, the setting of "source domain 100, target domain 1000" is not used, instead is "source domain 20, target domain 200".

The result as follows:

```
mnist->svhn: (not int paper)
```

```
Using downloaded and verified file: data/train_32x32.mat
Training start! | 0/1000 [00:00:00, -inf]
 0%|          | 0/1000 [35:33:296:58:46, 1071.27s/it]
 0%|          | 1/1000 [35:33:296:58:46, 1071.27s/it]
 0%|          | 2/1000 [35:33:296:58:46, 1071.27s/it]
 0%|          | 3/1000 [35:33:296:58:46, 1071.27s/it]
 0%|          | 4/1000 [35:33:296:58:46, 1071.27s/it]
 0%|          | 5/1000 [1:18:46:246:20:52, 891.31s/it]
 0%|          | 6/1000 [1:28:36:221:05:25, 800.73s/it]
 0%|          | 7/1000 [1:48:54:193:54:15, 703.58s/it]
 0%|          | 8/1000 [1:48:54:193:54:15, 703.58s/it]
 0%|          | 9/1000 [1:59:25:187:40:59, 661.80s/it]
 0%|          | 10/1000 [2:10:17:185:84:59, 673.03s/it]
 0%|          | 11/1000 [2:21:25:184:26:16, 671.36s/it]
 0%|          | 12/1000 [2:32:37:184:19:55, 671.46s/it]
 0%|          | 13/1000 [2:43:47:184:19:55, 671.46s/it]
 0%|          | 14/1000 [2:55:29:186:11:50, 679.83s/it]
 0%|          | 15/1000 [3:07:12:187:52:17, 606.65s/it]
 0%|          | 16/1000 [3:19:88:198:06:37, 655.53s/it]
 0%|          | 17/1000 [3:31:16:192:04:19, 705.26s/it]
 0%|          | 18/1000 [3:43:41:193:03:36, 717.18s/it]
 0%|          | 19/1000 [3:55:57:193:02:53, 738.11s/it]
 0%|          | 20/1000 [4:09:34:203:44:31, 748.44s/it]
 0%|          | 21/1000 [4:23:01:208:16:16, 765.86s/it]
 0%|          | 22/1000 [4:35:41:237:23:16, 800.24s/it]
 0%|          | 23/1000 [4:47:51:237:23:16, 837.51s/it]
 0%|          | 24/1000 [5:08:54:237:02:14, 874.32s/it]
 0%|          | 25/1000 [5:24:57:244:00:58, 900.38s/it]
 0%|          | 26/1000 [5:41:49:232:47:28, 934.34s/it]
 0%|          | 27/1000 [5:59:51:264:46:26, 978.48s/it]
 0%|          | 28/1000 [6:15:11:259:30:45, 961.16s/it]
 224it [01:10, 3.26it/s] Epoch 27 - S 99.9% - T 81.3%
```

In the `first iteration`, you can get more than `90%` in the training set and nearly `80%` in the test set. Obviously, it's little overfitting. But the speed and accuracy is good, respectively. Maybe in the setting of "source domain 100, target domain 1000", the result will be good.

Compare

Method	Domains (source → target)			
	MNIST → MNIST-M	Syn. Digits → SVHN	SVHN → MNIST	Syn. Signs → GTSRB
Transf. Repr. [22]	13.30	-	21.20	-
SA [8]	43.10	13.56	40.68	18.35
CORAL [24]	42.30	14.80	36.90	13.10
ADDA [30]	-	-	24.00	-
DANN [9]	23.33 (55.87 %)	8.91 (79.67 %)	26.15 (42.57 %)	11.35 (46.39 %)
DSN w/ DANN [3]	16.80 (63.18 %)	8.80 (78.95 %)	17.30 (58.31 %)	6.90 (54.42 %)
DSN w/ MMD [3]	19.50 (56.77 %)	11.50 (31.58 %)	27.80 (32.26 %)	7.40 (51.02 %)
MMD [15]	23.10	12.00	28.90	8.90
DA _{MMD}	22.90	19.14	28.48	10.69
Ours (DA _{assoc} fixed params [†])	10.47 ± 0.28	8.70 ± 0.2	4.32 ± 1.54	17.20 ± 1.32
Ours (DA _{assoc})	10.53 (85.94 %)	8.14 (87.78 %)	2.40 (93.71 %)	2.34 (81.23)
Source only	35.96	15.68	30.71	4.59
Target only	6.37	7.09	0.50	1.82

Since I used the data in `torchvision.datasets`, there are only svhn and mnist, so I didn't experiment with other datasets. It only needs to convert the data, but because of the long test time and big space of one data set in this experiment, so if time permits, I will add them.

MCD_UA

The author used a `three-layer fully-joined network` with a `batch size set to 32` and optimized the model `using a SGD` with a `learning rate of 1e-3`. And report accuracy after every ten iterations

The author gives the [MNIST data](#).

If you run an experiment on adaptation from svhn to mnist, where `num_k` indicates the number of update for generator.

```
python main.py --source svhn --target mnist --num_k 3
```

If you want to run an experiment **using gradient reversal layer**, simply add option `--one_step` when running this code:

```
python main.py --source svhn --target mnist --one_step
```

Since I integrated the code, please use `python main.py --model='MCD_UA'` and You can set other parameters, as explained above.

I chose two sets of experiments, one is `mnist to usps`; the other is `svhn to mnist`. And for the former to do a `three-layer fully connected network` and `using gradient reversal layer` comparison; for the latter to do a `three-layer and four-layer full-join network` comparison. Due to the high number of iterations, only the first and last two result images are released here. The complete results are in `results_image/MCD_UA_record/`.

Svhn to mnist

Three layer fully connected network

```

$ python main.py --source svhn --target mnist --num_k 3
Namespace(all_use='no', batch_size=128, checkpoint_dir='checkpoint', cuda=True, eval_onl
y=False, lr=0.0002, max_epoch=200, no_cuda=False, num_k=3, one_step=False, optimizer='ad
am', resume_epoch=100, save_epoch=10, save_model=False, seed=1, source='svhn', target='m
nist', use_abs_diff=False)
dataset loading
(55000, 3, 32, 32)
/usr/local/lib/python3.6/dist-packages/torchvision/transforms/transforms.py:208: UserWar
ning: The use of the transforms.Scale transform is deprecated, please use transforms.Res
ize instead.
    "please use transforms.Resize instead.")
load finished!
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:90: UserWarning: Implicit dimensio
n choice for softmax has been deprecated. Change the call to include dim=X as an argumen
t.
    return torch.mean(torch.abs(F.softmax(out1) - F.softmax(out2)))
Train Epoch: 0 [0/100 (0%)] Loss1: 1.911990 Loss2: 1.88967 Discrepancy: 0.021790
Train Epoch: 0 [100/100 (0%)] Loss1: 0.557463 Loss2: 0.56178 Discrepancy: 0.008060
Train Epoch: 0 [200/100 (0%)] Loss1: 0.406073 Loss2: 0.39611 Discrepancy: 0.008883
Train Epoch: 0 [300/100 (0%)] Loss1: 0.305152 Loss2: 0.29400 Discrepancy: 0.011216
Train Epoch: 0 [400/100 (1%)] Loss1: 0.346556 Loss2: 0.34085 Discrepancy: 0.009077
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:229: UserWarning: volatile was rem
oved and now has no effect. Use `with torch.no_grad():` instead.
    img, label = Variable(img, volatile=True), Variable(label)

Test set: Average loss: -0.0371, Accuracy C1: 6837/10000 (68%) Accuracy C2: 6838/10000 (
68%) Accuracy Ensemble: 6861/10000 (68%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 1 [0/100 (0%)] Loss1: 0.309628 Loss2: 0.30610 Discrepancy: 0.011315
Train Epoch: 1 [100/100 (0%)] Loss1: 0.303234 Loss2: 0.31203 Discrepancy: 0.012625
Train Epoch: 1 [200/100 (0%)] Loss1: 0.270616 Loss2: 0.28401 Discrepancy: 0.008394
Train Epoch: 1 [300/100 (0%)] Loss1: 0.375139 Loss2: 0.38299 Discrepancy: 0.009320
Train Epoch: 1 [400/100 (1%)] Loss1: 0.458935 Loss2: 0.45941 Discrepancy: 0.009006

Test set: Average loss: -0.0412, Accuracy C1: 7630/10000 (76%) Accuracy C2: 7587/10000 (
75%) Accuracy Ensemble: 7607/10000 (76%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 2 [0/100 (0%)] Loss1: 0.237049 Loss2: 0.24444 Discrepancy: 0.007628
Train Epoch: 2 [100/100 (0%)] Loss1: 0.331287 Loss2: 0.325654 Discrepancy: 0
.007541
Train Epoch: 2 [200/100 (0%)] Loss1: 0.459615 Loss2: 0.463207 Discrepancy: 0
.009048
Train Epoch: 2 [300/100 (0%)] Loss1: 0.256926 Loss2: 0.246781 Discrepancy: 0
.009858
Train Epoch: 2 [400/100 (1%)] Loss1: 0.259559 Loss2: 0.258995 Discrepancy: 0
.009695

Test set: Average loss: -0.0394, Accuracy C1: 7923/10000 (79%) Accuracy C2: 7901/10000 (
79%) Accuracy Ensemble: 7911/10000 (79%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 3 [0/100 (0%)] Loss1: 0.179793 Loss2: 0.182681 Discrepancy: 0
.007457
Train Epoch: 3 [100/100 (0%)] Loss1: 0.189167 Loss2: 0.186856 Discrepancy: 0
.008367
Train Epoch: 3 [200/100 (0%)] Loss1: 0.439096 Loss2: 0.449638 Discrepancy: 0
.008867

```

```

Train Epoch: 43 [200/100 (0%)] Loss1: 0.108076 Loss2: 0.105094 Disc
repancy: 0.001884
Train Epoch: 43 [300/100 (0%)] Loss1: 0.070828 Loss2: 0.069617 Disc
repancy: 0.001500
Train Epoch: 43 [400/100 (1%)] Loss1: 0.064212 Loss2: 0.066789 Disc
repancy: 0.001651

Test set: Average loss: -0.0773, Accuracy C1: 9591/10000 (95%) Accuracy C2: 95
93/10000 (95%) Accuracy Ensemble: 9588/10000 (95%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 44 [0/100 (0%)] Loss1: 0.032637 Loss2: 0.031198 Disc
repancy: 0.001405
Train Epoch: 44 [100/100 (0%)] Loss1: 0.047018 Loss2: 0.046000 Disc
repancy: 0.002637
Train Epoch: 44 [200/100 (0%)] Loss1: 0.083661 Loss2: 0.084633 Disc
repancy: 0.002645
Train Epoch: 44 [300/100 (0%)] Loss1: 0.098382 Loss2: 0.095631 Disc
repancy: 0.002202
Train Epoch: 44 [400/100 (1%)] Loss1: 0.071495 Loss2: 0.075741 Disc
repancy: 0.002169

Test set: Average loss: -0.0796, Accuracy C1: 9592/10000 (95%) Accuracy C2: 95
77/10000 (95%) Accuracy Ensemble: 9585/10000 (95%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 45 [0/100 (0%)] Loss1: 0.052420 Loss2: 0.049829 Disc
repancy: 0.001848
Train Epoch: 45 [100/100 (0%)] Loss1: 0.121637 Loss2: 0.124418 Disc
repancy: 0.001917
Train Epoch: 45 [200/100 (0%)] Loss1: 0.075262 Loss2: 0.074892 Disc
repancy: 0.002314
Train Epoch: 45 [300/100 (0%)] Loss1: 0.065439 Loss2: 0.062935 Disc
repancy: 0.001621
Train Epoch: 45 [400/100 (1%)] Loss1: 0.079050 Loss2: 0.080480 Disc
repancy: 0.001592

Test set: Average loss: -0.0739, Accuracy C1: 9645/10000 (96%) Accuracy C2: 96
43/10000 (96%) Accuracy Ensemble: 9644/10000 (96%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
Train Epoch: 46 [0/100 (0%)] Loss1: 0.067377 Loss2: 0.068424 Disc
repancy: 0.001051
Train Epoch: 46 [100/100 (0%)] Loss1: 0.104238 Loss2: 0.099360 Disc
repancy: 0.001753
Train Epoch: 46 [200/100 (0%)] Loss1: 0.058262 Loss2: 0.057898 Disc
repancy: 0.001712
Train Epoch: 46 [300/100 (0%)] Loss1: 0.086384 Loss2: 0.084947 Disc
repancy: 0.001438
Train Epoch: 46 [400/100 (1%)] Loss1: 0.065899 Loss2: 0.067052 Disc
repancy: 0.001809

Test set: Average loss: -0.0785, Accuracy C1: 9624/10000 (96%) Accuracy C2: 96
50/10000 (96%) Accuracy Ensemble: 9640/10000 (96%)

recording %s record/svhn_mnist_k_3_onestep_False_1_test.txt
root at 641aa73707ea in /home/huangchenxi/transfer-learning/MCD_UDA
$ 

```

using gradient reversal layer

```

$ python main.py --source svhn --target mnist --one_step
Namespace(all_use='no', batch_size=128, checkpoint_dir='checkpoint', cuda=True, eval_only=False, lr=0.0002, max_epoch=200, no_cuda=False, num_k=4, one_step=True, optimizer='adam', resume_epoch=100, save_epoch=10, save_model=False, seed=1, source='svhn', target='mnist', use_abs_diff=False)
dataset loading
(55000, 3, 32, 32)
/usr/local/lib/python3.6/dist-packages/torchvision/transforms/transforms.py:208: UserWarning: The use of the transforms.Scale transform is deprecated, please use transforms.Resize instead.
  "please use transforms.Resize instead.")
load finished!
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:90: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.
  return torch.mean(torch.abs(F.softmax(out1) - F.softmax(out2)))
Train Epoch: 0 [0/100 (0%)] Loss1: 2.268277 Loss2: 2.305896 Discrepancy: -0.024503
Train Epoch: 0 [100/100 (0%)] Loss1: 0.713227 Loss2: 0.706736 Discrepancy: -0.006221
Train Epoch: 0 [200/100 (0%)] Loss1: 0.387395 Loss2: 0.38860 Discrepancy: -0.006290
Train Epoch: 0 [300/100 (0%)] Loss1: 0.397803 Loss2: 0.397506 Discrepancy: -0.005612
Train Epoch: 0 [400/100 (1%)] Loss1: 0.365004 Loss2: 0.369847 Discrepancy: -0.004369
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:229: UserWarning: volatile was removed and now has no effect. Use `with torch.no_grad():` instead.
  img, label = Variable(img, volatile=True), Variable(label)

Test set: Average loss: -0.0269, Accuracy C1: 5740/10000 (57%) Accuracy C2: 5749/10000 (57%) Accuracy Ensemble: 5746/10000 (57%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 1 [0/100 (0%)] Loss1: 0.404227 Loss2: 0.402911 Discrepancy: -0.005601
Train Epoch: 1 [100/100 (0%)] Loss1: 0.322874 Loss2: 0.324709 Discrepancy: -0.004417
Train Epoch: 1 [200/100 (0%)] Loss1: 0.313822 Loss2: 0.311150 Discrepancy: -0.004363
Train Epoch: 1 [300/100 (0%)] Loss1: 0.391723 Loss2: 0.391865 Discrepancy: -0.004365
Train Epoch: 1 [400/100 (1%)] Loss1: 0.444960 Loss2: 0.441041 Discrepancy: -0.004264

Test set: Average loss: -0.0313, Accuracy C1: 6472/10000 (64%) Accuracy C2: 6459/10000 (64%) Accuracy Ensemble: 6467/10000 (64%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 2 [0/100 (0%)] Loss1: 0.235934 Loss2: 0.235478 Discrepancy: -0.004050
Train Epoch: 2 [100/100 (0%)] Loss1: 0.288127 Loss2: 0.281232 Discrepancy: -0.003681
Train Epoch: 2 [200/100 (0%)] Loss1: 0.517296 Loss2: 0.519623 Discrepancy: -0.004146
Train Epoch: 2 [300/100 (0%)] Loss1: 0.278004 Loss2: 0.281306 Discrepancy: -0.002996
Train Epoch: 2 [400/100 (1%)] Loss1: 0.269612 Loss2: 0.268115 Discrepancy: -0.003091

Test set: Average loss: -0.0313, Accuracy C1: 6145/10000 (61%) Accuracy C2: 6140/10000 (61%) Accuracy Ensemble: 6139/10000 (61%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 3 [0/100 (0%)] Loss1: 0.275257 Loss2: 0.278027 Discrepancy: -0.003090
Train Epoch: 3 [100/100 (0%)] Loss1: 0.162313 Loss2: 0.160310 Discrepancy: -0.004580
Train Epoch: 3 [200/100 (0%)] Loss1: 0.364433 Loss2: 0.365316 Discrepancy: -0.002920
Train Epoch: 3 [300/100 (0%)] Loss1: 0.350219 Loss2: 0.351968 Discrepancy: -0.002836
Train Epoch: 3 [400/100 (1%)] Loss1: 0.125356 Loss2: 0.126044 Discrepancy: -0.003095

Test set: Average loss: -0.0339, Accuracy C1: 6388/10000 (63%) Accuracy C2: 6357/10000 (63%) Accuracy Ensemble: 6372/10000 (63%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 4 [0/100 (0%)] Loss1: 0.186891 Loss2: 0.184971 Discrepancy: -0.003328
Train Epoch: 4 [100/100 (0%)] Loss1: 0.232140 Loss2: 0.236099 Discrepancy: -0.002692
Train Epoch: 4 [200/100 (0%)] Loss1: 0.248813 Loss2: 0.246535 Discrepancy: -0.002602
Train Epoch: 4 [300/100 (0%)] Loss1: 0.247648 Loss2: 0.244619 Discrepancy: -0.003100
Train Epoch: 4 [400/100 (1%)] Loss1: 0.265795 Loss2: 0.269968 Discrepancy: -0.002684

Test set: Average loss: -0.0319, Accuracy C1: 6604/10000 (66%) Accuracy C2: 6607/10000 (66%) Accuracy Ensemble: 6606/10000 (66%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 5 [0/100 (0%)] Loss1: 0.315748 Loss2: 0.317907 Discrepancy: -0.002660
Train Epoch: 5 [100/100 (0%)] Loss1: 0.180223 Loss2: 0.184398 Discrepancy: -0.002916

```

```

Test set: Average loss: -0.0433, Accuracy C1: 6629/10000 (66%) Accuracy C2: 6628/10000 (66%) Accuracy Ensemble: 6623/10000 (66%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 41 [0/100 (0%)] Loss1: 0.108207 Loss2: 0.108199 Discrepancy: -0.001499
Train Epoch: 41 [100/100 (0%)] Loss1: 0.016104 Loss2: 0.016335 Discrepancy: -0.001511
Train Epoch: 41 [200/100 (0%)] Loss1: 0.032615 Loss2: 0.033705 Discrepancy: -0.001385
Train Epoch: 41 [300/100 (0%)] Loss1: 0.025081 Loss2: 0.025062 Discrepancy: -0.001392
Train Epoch: 41 [400/100 (1%)] Loss1: 0.051415 Loss2: 0.052083 Discrepancy: -0.001751

Test set: Average loss: -0.0443, Accuracy C1: 6806/10000 (68%) Accuracy C2: 6801/10000 (68%) Accuracy Ensemble: 6803/10000 (68%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 42 [0/100 (0%)] Loss1: 0.015099 Loss2: 0.015437 Discrepancy: -0.001461
Train Epoch: 42 [100/100 (0%)] Loss1: 0.048697 Loss2: 0.048781 Discrepancy: -0.001620
Train Epoch: 42 [200/100 (0%)] Loss1: 0.032834 Loss2: 0.031839 Discrepancy: -0.001269
Train Epoch: 42 [300/100 (0%)] Loss1: 0.037795 Loss2: 0.038357 Discrepancy: -0.001927
Train Epoch: 42 [400/100 (1%)] Loss1: 0.038690 Loss2: 0.039036 Discrepancy: -0.001380

Test set: Average loss: -0.0498, Accuracy C1: 7128/10000 (71%) Accuracy C2: 7140/10000 (71%) Accuracy Ensemble: 7136/10000 (71%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 43 [0/100 (0%)] Loss1: 0.111399 Loss2: 0.112127 Discrepancy: -0.001343
Train Epoch: 43 [100/100 (0%)] Loss1: 0.039488 Loss2: 0.038493 Discrepancy: -0.001622
Train Epoch: 43 [200/100 (0%)] Loss1: 0.049160 Loss2: 0.050125 Discrepancy: -0.001262
Train Epoch: 43 [300/100 (0%)] Loss1: 0.026048 Loss2: 0.025455 Discrepancy: -0.001060
Train Epoch: 43 [400/100 (1%)] Loss1: 0.066523 Loss2: 0.066126 Discrepancy: -0.001474

Test set: Average loss: -0.0474, Accuracy C1: 6965/10000 (69%) Accuracy C2: 6962/10000 (69%) Accuracy Ensemble: 6965/10000 (69%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 44 [0/100 (0%)] Loss1: 0.024393 Loss2: 0.024379 Discrepancy: -0.001475
Train Epoch: 44 [100/100 (0%)] Loss1: 0.031503 Loss2: 0.031038 Discrepancy: -0.001384
Train Epoch: 44 [200/100 (0%)] Loss1: 0.053952 Loss2: 0.054666 Discrepancy: -0.001014
Train Epoch: 44 [300/100 (0%)] Loss1: 0.090659 Loss2: 0.091470 Discrepancy: -0.001779
Train Epoch: 44 [400/100 (1%)] Loss1: 0.056462 Loss2: 0.058524 Discrepancy: -0.001342

Test set: Average loss: -0.0481, Accuracy C1: 6966/10000 (69%) Accuracy C2: 6966/10000 (69%) Accuracy Ensemble: 6969/10000 (69%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 45 [0/100 (0%)] Loss1: 0.031661 Loss2: 0.032142 Discrepancy: -0.001419
Train Epoch: 45 [100/100 (0%)] Loss1: 0.081738 Loss2: 0.082430 Discrepancy: -0.001402
Train Epoch: 45 [200/100 (0%)] Loss1: 0.050895 Loss2: 0.051133 Discrepancy: -0.001484
Train Epoch: 45 [300/100 (0%)] Loss1: 0.070765 Loss2: 0.070488 Discrepancy: -0.001254
Train Epoch: 45 [400/100 (1%)] Loss1: 0.066423 Loss2: 0.065869 Discrepancy: -0.001229

Test set: Average loss: -0.0494, Accuracy C1: 6863/10000 (68%) Accuracy C2: 6853/10000 (68%) Accuracy Ensemble: 6854/10000 (68%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
Train Epoch: 46 [0/100 (0%)] Loss1: 0.057762 Loss2: 0.058632 Discrepancy: -0.001192
Train Epoch: 46 [100/100 (0%)] Loss1: 0.028099 Loss2: 0.028188 Discrepancy: -0.001630
Train Epoch: 46 [200/100 (0%)] Loss1: 0.025466 Loss2: 0.025755 Discrepancy: -0.001336
Train Epoch: 46 [300/100 (0%)] Loss1: 0.033726 Loss2: 0.033338 Discrepancy: -0.001717
Train Epoch: 46 [400/100 (1%)] Loss1: 0.053387 Loss2: 0.053617 Discrepancy: -0.001337

Test set: Average loss: -0.0467, Accuracy C1: 6762/10000 (67%) Accuracy C2: 6758/10000 (67%) Accuracy Ensemble: 6759/10000 (67%)

recording %s record/svhn_mnist_k_4_onestep_True_0_test.txt
root at 60a5497e4409 in /home/huangchenxi/transfer-learning/MCD_UDA
$ []

```

Mnist to usps

Three layer fully connected network

```
$ python main.py --source mnist --target usps --num_k 3
Namespace(all_use='no', batch_size=128, checkpoint_dir='checkpoint', cuda=True, eval_only=False, lr=0.0002, max_epoch=200, no_cuda=False, num_k=3, one_step=False, optimizer='adam', resume_epoch=100, save_epoch=10, save_mod el=False, seed=1, source='mnist', target='usps', use_abs_diff=False)
dataset loading
(2000, 1, 28, 28)
/usr/local/lib/python3.6/dist-packages/torchvision/transforms/transforms.py:208: UserWarning: The use of the tr ansforms.Scale transform is deprecated, please use transforms.Resize instead.
  "please use transforms.Resize instead.")
load finished!
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:90: UserWarning: Implicit dimension choice for softmax ha s been deprecated. Change the call to include dim=X as an argument.
  return torch.mean(torch.abs(F.softmax(out1) - F.softmax(out2)))
Train Epoch: 0 [0/100 (0%)]    Loss1: 2.382223  Loss2: 2.382314      Discrepancy: 0.036035
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:229: UserWarning: volatile was removed and now has no eff ect. Use `with torch.no_grad():` instead.
  img, label = Variable(img, volatile=True), Variable(label)

Test set: Average loss: -0.0002, Accuracy C1: 712/1860 (38%) Accuracy C2: 491/1860 (26%) Accuracy Ensemble: 692 /1860 (37%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 1 [0/100 (0%)]    Loss1: 2.216476  Loss2: 2.210045      Discrepancy: 0.034149

Test set: Average loss: -0.0011, Accuracy C1: 968/1860 (52%) Accuracy C2: 693/1860 (37%) Accuracy Ensemble: 829 /1860 (44%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 2 [0/100 (0%)]    Loss1: 2.124503  Loss2: 2.066645      Discrepancy: 0.033762

Test set: Average loss: -0.0017, Accuracy C1: 993/1860 (53%) Accuracy C2: 785/1860 (42%) Accuracy Ensemble: 889 /1860 (47%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 3 [0/100 (0%)]    Loss1: 1.975672  Loss2: 1.926368      Discrepancy: 0.037194

Test set: Average loss: -0.0025, Accuracy C1: 951/1860 (51%) Accuracy C2: 823/1860 (44%) Accuracy Ensemble: 887 /1860 (47%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 4 [0/100 (0%)]    Loss1: 1.919788  Loss2: 1.755128      Discrepancy: 0.039241

Test set: Average loss: -0.0032, Accuracy C1: 988/1860 (53%) Accuracy C2: 924/1860 (49%) Accuracy Ensemble: 962 /1860 (51%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 5 [0/100 (0%)]    Loss1: 1.657431  Loss2: 1.705456      Discrepancy: 0.040850

Test set: Average loss: -0.0039, Accuracy C1: 1028/1860 (55%) Accuracy C2: 938/1860 (50%) Accuracy Ensemble: 98 8/1860 (53%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 6 [0/100 (0%)]    Loss1: 1.614554  Loss2: 1.560779      Discrepancy: 0.043945

Test set: Average loss: -0.0047, Accuracy C1: 1105/1860 (59%) Accuracy C2: 1037/1860 (55%) Accuracy Ensemble: 1 089/1860 (58%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 7 [0/100 (0%)]    Loss1: 1.541687  Loss2: 1.541064      Discrepancy: 0.048636

Test set: Average loss: -0.0052, Accuracy C1: 1054/1860 (56%) Accuracy C2: 1015/1860 (54%) Accuracy Ensemble: 1 037/1860 (55%)

recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt
Train Epoch: 8 [0/100 (0%)]    Loss1: 1.494078  Loss2: 1.384127      Discrepancy: 0.050657

Test set: Average loss: -0.0059, Accuracy C1: 1095/1860 (58%) Accuracy C2: 1077/1860 (57%) Accuracy Ensemble: 1 098/1860 (59%)
```

```
Test set: Average loss: -0.0492, Accuracy C1: 1762/1860 (94%) Accuracy C2: 1778/1860 (95%) Accuracy Ensemble: 1  
768/1860 (95%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 190 [0/100 (0%)] Loss1: 0.011243 Loss2: 0.010493 Discrepancy: 0.011267  
  
Test set: Average loss: -0.0491, Accuracy C1: 1757/1860 (94%) Accuracy C2: 1774/1860 (95%) Accuracy Ensemble: 1  
765/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 191 [0/100 (0%)] Loss1: 0.014260 Loss2: 0.018318 Discrepancy: 0.017557  
  
Test set: Average loss: -0.0497, Accuracy C1: 1746/1860 (93%) Accuracy C2: 1768/1860 (95%) Accuracy Ensemble: 1  
761/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 192 [0/100 (0%)] Loss1: 0.007250 Loss2: 0.031433 Discrepancy: 0.012802  
  
Test set: Average loss: -0.0484, Accuracy C1: 1750/1860 (94%) Accuracy C2: 1764/1860 (94%) Accuracy Ensemble: 1  
757/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 193 [0/100 (0%)] Loss1: 0.011271 Loss2: 0.007898 Discrepancy: 0.021786  
  
Test set: Average loss: -0.0490, Accuracy C1: 1745/1860 (93%) Accuracy C2: 1769/1860 (95%) Accuracy Ensemble: 1  
765/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 194 [0/100 (0%)] Loss1: 0.014166 Loss2: 0.013017 Discrepancy: 0.022087  
  
Test set: Average loss: -0.0483, Accuracy C1: 1762/1860 (94%) Accuracy C2: 1780/1860 (95%) Accuracy Ensemble: 1  
770/1860 (95%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 195 [0/100 (0%)] Loss1: 0.024037 Loss2: 0.013002 Discrepancy: 0.014850  
  
Test set: Average loss: -0.0508, Accuracy C1: 1766/1860 (94%) Accuracy C2: 1786/1860 (96%) Accuracy Ensemble: 1  
781/1860 (95%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 196 [0/100 (0%)] Loss1: 0.020742 Loss2: 0.007589 Discrepancy: 0.018010  
  
Test set: Average loss: -0.0491, Accuracy C1: 1750/1860 (94%) Accuracy C2: 1767/1860 (95%) Accuracy Ensemble: 1  
765/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 197 [0/100 (0%)] Loss1: 0.013722 Loss2: 0.006397 Discrepancy: 0.014234  
  
Test set: Average loss: -0.0498, Accuracy C1: 1753/1860 (94%) Accuracy C2: 1775/1860 (95%) Accuracy Ensemble: 1  
763/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 198 [0/100 (0%)] Loss1: 0.011574 Loss2: 0.016353 Discrepancy: 0.012327  
  
Test set: Average loss: -0.0480, Accuracy C1: 1742/1860 (93%) Accuracy C2: 1763/1860 (94%) Accuracy Ensemble: 1  
759/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
Train Epoch: 199 [0/100 (0%)] Loss1: 0.008313 Loss2: 0.015026 Discrepancy: 0.013231  
  
Test set: Average loss: -0.0473, Accuracy C1: 1747/1860 (93%) Accuracy C2: 1765/1860 (94%) Accuracy Ensemble: 1  
755/1860 (94%)  
  
recording %s record/mnist_usps_k_3_alluse_no_onestep_False_0_test.txt  
root at 263b4efd5a52 in /home/huangchenxi/transfer-learning/MCD_UDA  
$ 
```

Four layer fully connected network

```
[\$ python main.py --source mnist --target usps --num_k 4
Namespace(all_use='no', batch_size=128, checkpoint_dir='checkpoint', cuda=True, eval_only=False, lr=0.0002, max_
_epoch=200, no_cuda=False, num_k=4, one_step=False, optimizer='adam', resume_epoch=100, save_epoch=10, save_mod_
el=False, seed=1, source='mnist', target='usps', use_abs_diff=False)
dataset loading
(2000, 1, 28, 28)
/usr/local/lib/python3.6/dist-packages/torchvision/transforms/transforms.py:208: UserWarning: The use of the tr_
ansforms.Scale transform is deprecated, please use transforms.Resize instead.
  "please use transforms.Resize instead.")
load finished!
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:90: UserWarning: Implicit dimension choice for softmax ha_
s been deprecated. Change the call to include dim=X as an argument.
  return torch.mean(torch.abs(F.softmax(out1) - F.softmax(out2)))
Train Epoch: 0 [0/100 (0%)]    Loss1: 2.343074    Loss2: 2.354578      Discrepancy: 0.034134
/home/huangchenxi/transfer-learning/MCD_UDA/solver.py:229: UserWarning: volatile was removed and now has no eff_
ect. Use `with torch.no_grad():` instead.
  img, label = Variable(img, volatile=True), Variable(label)

Test set: Average loss: -0.0003, Accuracy C1: 641/1860 (34%) Accuracy C2: 490/1860 (26%) Accuracy Ensemble: 719
/1860 (38%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 1 [0/100 (0%)]    Loss1: 2.217272    Loss2: 2.172659      Discrepancy: 0.033076

Test set: Average loss: -0.0011, Accuracy C1: 988/1860 (53%) Accuracy C2: 758/1860 (40%) Accuracy Ensemble: 915
/1860 (49%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 2 [0/100 (0%)]    Loss1: 2.168787    Loss2: 2.101144      Discrepancy: 0.033594

Test set: Average loss: -0.0019, Accuracy C1: 1033/1860 (55%) Accuracy C2: 889/1860 (47%) Accuracy Ensemble: 98
7/1860 (53%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 3 [0/100 (0%)]    Loss1: 1.990340    Loss2: 2.051481      Discrepancy: 0.035275

Test set: Average loss: -0.0027, Accuracy C1: 1093/1860 (58%) Accuracy C2: 982/1860 (52%) Accuracy Ensemble: 10
65/1860 (57%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 4 [0/100 (0%)]    Loss1: 1.879270    Loss2: 1.876694      Discrepancy: 0.036969

Test set: Average loss: -0.0035, Accuracy C1: 1062/1860 (57%) Accuracy C2: 1005/1860 (54%) Accuracy Ensemble: 1
056/1860 (56%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 5 [0/100 (0%)]    Loss1: 1.660003    Loss2: 1.677380      Discrepancy: 0.039429

Test set: Average loss: -0.0044, Accuracy C1: 1131/1860 (60%) Accuracy C2: 1094/1860 (58%) Accuracy Ensemble: 1
141/1860 (61%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 6 [0/100 (0%)]    Loss1: 1.609736    Loss2: 1.597398      Discrepancy: 0.042612

Test set: Average loss: -0.0052, Accuracy C1: 1160/1860 (62%) Accuracy C2: 1131/1860 (60%) Accuracy Ensemble: 1
163/1860 (62%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 7 [0/100 (0%)]    Loss1: 1.582161    Loss2: 1.559811      Discrepancy: 0.047276

Test set: Average loss: -0.0060, Accuracy C1: 1222/1860 (65%) Accuracy C2: 1201/1860 (64%) Accuracy Ensemble: 1
236/1860 (66%)

recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt
Train Epoch: 8 [0/100 (0%)]    Loss1: 1.454922    Loss2: 1.429422      Discrepancy: 0.048694

Test set: Average loss: -0.0070, Accuracy C1: 1251/1860 (67%) Accuracy C2: 1221/1860 (65%) Accuracy Ensemble: 1
254/1860 (67%)
```

```
Test set: Average loss: -0.0504, Accuracy C1: 1764/1860 (94%) Accuracy C2: 1758/1860 (94%) Accuracy Ensemble: 1  
762/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 190 [0/100 (0%)] Loss1: 0.022510 Loss2: 0.009406 Discrepancy: 0.011573  
  
Test set: Average loss: -0.0523, Accuracy C1: 1772/1860 (95%) Accuracy C2: 1757/1860 (94%) Accuracy Ensemble: 1  
763/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 191 [0/100 (0%)] Loss1: 0.015358 Loss2: 0.005861 Discrepancy: 0.012024  
  
Test set: Average loss: -0.0498, Accuracy C1: 1761/1860 (94%) Accuracy C2: 1753/1860 (94%) Accuracy Ensemble: 1  
760/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 192 [0/100 (0%)] Loss1: 0.046938 Loss2: 0.011934 Discrepancy: 0.009163  
  
Test set: Average loss: -0.0497, Accuracy C1: 1756/1860 (94%) Accuracy C2: 1733/1860 (93%) Accuracy Ensemble: 1  
746/1860 (93%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 193 [0/100 (0%)] Loss1: 0.012432 Loss2: 0.011339 Discrepancy: 0.014111  
  
Test set: Average loss: -0.0510, Accuracy C1: 1776/1860 (95%) Accuracy C2: 1762/1860 (94%) Accuracy Ensemble: 1  
770/1860 (95%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 194 [0/100 (0%)] Loss1: 0.011661 Loss2: 0.008523 Discrepancy: 0.013203  
  
Test set: Average loss: -0.0503, Accuracy C1: 1768/1860 (95%) Accuracy C2: 1758/1860 (94%) Accuracy Ensemble: 1  
765/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 195 [0/100 (0%)] Loss1: 0.010646 Loss2: 0.021519 Discrepancy: 0.013082  
  
Test set: Average loss: -0.0501, Accuracy C1: 1768/1860 (95%) Accuracy C2: 1751/1860 (94%) Accuracy Ensemble: 1  
761/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 196 [0/100 (0%)] Loss1: 0.018338 Loss2: 0.008945 Discrepancy: 0.010631  
  
Test set: Average loss: -0.0515, Accuracy C1: 1786/1860 (96%) Accuracy C2: 1774/1860 (95%) Accuracy Ensemble: 1  
777/1860 (95%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 197 [0/100 (0%)] Loss1: 0.010997 Loss2: 0.012398 Discrepancy: 0.008331  
  
Test set: Average loss: -0.0487, Accuracy C1: 1764/1860 (94%) Accuracy C2: 1744/1860 (93%) Accuracy Ensemble: 1  
754/1860 (94%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
Train Epoch: 198 [0/100 (0%)] Loss1: 0.014923 Loss2: 0.025930 Discrepancy: 0.014242  
  
Test set: Average loss: -0.0509, Accuracy C1: 1771/1860 (95%) Accuracy C2: 1762/1860 (94%) Accuracy Ensemble: 1  
773/1860 (95%)  
  
recording %s record/mnist_usps_k_4_alluse_no_onestep_False_0_test.txt  
  
root at 263b4efd5a52 in /home/huangchenxi/transfer-learning/MCD_UDA  
$ 
```

Compare

ACC	M2U-3	M2U-4	S2M-3	S2M-3-one-step
In my test	94.35%	95.65%	96.4%	67.59%
In paper	93.8 ± 0.8	94.2 ± 0.7	95.9 ± 0.5	

The result **is consistent with** that in paper, which also shows that the MCD_UDA algorithm is indeed more accurate. Using one step, the accuracy is also the expected to be poor.

MADA

Because there only codes in caffe, that it is difficult to recurrent it. So first I read paper again, then I write the network of MADA. After that, I use the architecture that I used before to integrate it.

The network architecture as follows:

```
def __init__(self, n_classes, convnet=None, classifier=None):
    super().__init__()

    self._n_classes = n_classes

    self._convnet = convnet or ConvNet()
    self._classifier = classifier or Classifier(n_classes, 12544)
    self._grl = GRL(factor=-1)
    self._domain_classifiers = [  # k classifiers
        Classifier(1, 12544)
        for _ in range(n_classes)
    ]
```

Compare

The result shows the algorithm is excellent. Because I don't have much time to do the remain work, so I don't compare it.

```
root at 61cebc9a3a1 in /home/huangchenxi/transfer-learning/transfer-learning
[$ python main.py --model='MADA' --source='amazon' --target='dslr' --gpu=1 --num_workers=4 --num_workers=0 --batch_size=32
MADA
Data size = 2817 , corrects = 3
Initial Train Loss: 0.0000 Acc: 0.0011
Data size = 498 , corrects = 2
Initial Test Loss: 0.0000 Acc: 0.0011

Epoch 0/255
iteration : 0
Data size = 512 , corrects = 480
Alpha = 0.0

Train Loss: 21.7576 Acc: 0.9375
Data size = 498 , corrects = 498
Test Loss: 0.0000 Acc: 1.0000
Current Best Test Acc : 0.937500    Current Best Test Loss : 0.000000  Cur lr : 0
.000300

Epoch 1/255
iteration : 16
Data size = 512 , corrects = 512
Alpha = 0.019528766852031865

Train Loss: 21.4984 Acc: 1.0000
Data size = 498 , corrects = 498
Test Loss: 0.0000 Acc: 1.0000
Current Best Test Acc : 1.000000    Current Best Test Loss : 0.000000  Cur lr : 0
.000292

Epoch 2/255
iteration : 32
Data size = 512 , corrects = 512
Alpha = 0.03904264390418577

Train Loss: 21.4917 Acc: 1.0000
Data size = 498 , corrects = 498
Test Loss: 0.0000 Acc: 1.0000
Current Best Test Acc : 1.000000    Current Best Test Loss : 0.000000  Cur lr : 0
.000284

Epoch 3/255
iteration : 48
Data size = 512 , corrects = 512
Alpha = 0.05852678673317713

Train Loss: 21.4991 Acc: 1.0000
Data size = 498 , corrects = 498
Test Loss: 0.0000 Acc: 1.0000
Current Best Test Acc : 1.000000    Current Best Test Loss : 0.000000  Cur lr : 0
.000276

Epoch 4/255
iteration : 64

Data size = 512 , corrects = 512
Alpha = 0.07796644137536823
```

Of course, this is a bit strange. In the second epoch, the accuracy rate has been 100%. In the paper, the effect of A->D is not so good, although in W->D, it can indeed reach 100%.

Table 1: Accuracy (%) on *Office-31* for unsupervised domain adaptation (AlexNet and ResNet)

Method	A → W	D → W	W → D	A → D	D → A	W → A	Avg
AlexNet (Krizhevsky, Sutskever, and Hinton 2012)	60.6±0.4	95.4±0.2	99.0±0.1	64.2±0.3	45.5±0.5	48.3±0.5	68.8
	59.0±0.0	90.2±0.0	88.2±0.0	57.8±0.0	51.6±0.0	47.9±0.0	65.8
	58.4±0.0	93.6±0.0	91.0±0.0	58.6±0.0	52.4±0.0	46.1±0.0	66.7
	61.0±0.5	95.0±0.3	98.5±0.3	64.9±0.4	47.2±0.5	49.4±0.4	69.3
	68.5±0.3	96.0±0.1	99.0±0.1	66.8±0.2	50.0±0.4	49.8±0.3	71.7
	73.3±0.2	96.8±0.2	99.6±0.1	71.0±0.2	50.5±0.3	51.0±0.1	73.7
	73.0±0.5	96.4±0.3	99.2±0.3	72.3±0.3	52.4±0.4	50.4±0.5	74.1
RevGrad (Ganin and Lempitsky 2015)							
MADA	78.5±0.2	99.8±0.1	100.0±0.0	74.1±0.1	56.0±0.2	54.5±0.3	77.1
ResNet (He et al. 2016)	68.4±0.2	96.7±0.1	99.3±0.1	68.9±0.2	62.5±0.3	60.7±0.3	76.1
	74.7±0.0	96.7±0.0	99.6±0.0	76.1±0.0	63.7±0.0	62.9±0.0	79.3
	74.8±0.0	95.0±0.0	98.2±0.0	76.5±0.0	65.4±0.0	63.0±0.0	78.8
	75.8±0.2	95.0±0.2	98.2±0.1	77.5±0.3	67.4±0.4	64.0±0.5	79.7
	83.8±0.4	96.8±0.2	99.5±0.1	78.4±0.2	66.7±0.3	62.7±0.2	81.3
	84.5±0.2	96.8±0.1	99.4±0.1	77.5±0.3	66.2±0.2	64.8±0.3	81.6
	82.0±0.4	96.9±0.2	99.1±0.1	79.7±0.4	68.2±0.4	67.4±0.5	82.2
MADA	90.0±0.1	97.4±0.1	99.6±0.1	87.8±0.2	70.3±0.3	66.4±0.3	85.2

Post-preface

Every experiment has been reproduced. Since I first learned about migration and learning, the code may be a little sloppy. If you have some problems, please contact to me.